

CS2 AI COACHING SYSTEM

Ultimate CS2 Coach

Part 3 — Application, UI, Tools and Build

Application Logic, UI Qt/PySide6, Ingestion Pipeline, 17 Diagnostic Tools,
81 Test Files, Remediation

Autore: Renan Augusto Macena

Marzo 2026

Ultimate CS2 Coach — Part 3: Program, UI, Tools, and Build

Topics: Complete database schema (three-tier storage, SQLite WAL, SQLAlchemy ORM), Training regime and maturity gates (CALIBRATING/LEARNING/MATURE), Loss function catalog, Complete Program Logic (from launch to advice): Session Engine, Digester, Teacher, Hunter, Pulse; Qt/PySide6 desktop interface (13 screens, ViewModels, Qt Signals); Ingestion pipeline (demo and pro); Unified control console; New user onboarding; Storage architecture; Playback engine and tactical viewer; Spatial data and maps; Observability and logging; Reporting; Quotas and limits; Fault tolerance; Complete user journey (4 flows); Tool suite (validation and diagnostics); Test suite; Pre-commit hooks; Build, packaging, deployment; Alembic migrations; HLTV sync service; RASP Guard; MatchVisualizer; Runtime configuration files; Root-level entry points; Glossary.

Author: Renan Augusto Macena

This document is the continuation of [Part 2 — Services, Analysis, and Database](#), which covers coaching services, analysis engines, knowledge, and the Database. Part 3 drops down to the program level: how the entire system boots, how the daemons orchestrate, how the desktop interface makes coaching visible, how the code is tested, validated, and deployed.

Table of Contents

Part 3 — Program, UI, Tools, and Build (this document)

9. [Database schema and data lifecycle](#)
10. [Training regime and maturity gates](#)
11. [Loss function catalog](#)
12. [Complete Program Logic — From Launch to Advice](#)
 - 12.1 Application startup and bootstrap
 - 12.2 Session Engine and Quad-Daemon
 - 12.3 Digester Daemon and ingestion pipeline
 - 12.4 Teacher Daemon and training orchestrator
 - 12.5 Desktop Interface (Qt/PySide6, 13 screens, ViewModels)
 - 12.6 Ingestion Pipeline ([ingestion/](#))

- 12.7 Unified Control Console ([backend/control/](#))
- 12.8 Onboarding and New User Flow
- 12.9 Storage Architecture ([backend/storage/](#))
- 12.10 Playback Engine and Tactical Viewer
- 12.11 Spatial Data and Map Management
- 12.12 Observability and Logging
- 12.13 Reporting and Visualization
- 12.14 Quota and Limit Management
- 12.15 Fault Tolerance and Recovery
- 12.16 Complete User Journey — 4 Main Flows
- 12.17 Tool Suite — Validation and Diagnostics ([tools/](#))
- 12.18 Test Suite Architecture ([tests/](#))
- 12.19 The 12 Phases of Systematic Remediation
- 12.20 Pre-commit Hooks and Quality Gates
- 12.21 Build, Packaging, and Deployment
- 12.22 Alembic Migration System
- 12.23 Main Ingestion Orchestrator ([run_ingestion.py](#))
- 12.24 HLTV Sync Service and Background Daemon
- 12.25 RASP Guard — Runtime Code Integrity
- 12.26 MatchVisualizer — Advanced Rendering
- 12.27 Runtime Data and Configuration Files
- 12.28 Root-Level Entry Points
- [Architectural Summary](#)
- [Cross-Part Interconnection Map](#)
- [Technical Glossary](#)

9. Database schema and data lifecycle

The project uses **SQLModel** (Pydantic + SQLAlchemy) with SQLite (WAL mode) and a specialized **three-tier storage architecture**. In total **24 SQLModel tables** ([backend/storage/db_models.py](#) : 21 monolith+HLTV tables; [backend/storage/match_data_manager.py](#) : 3 per-match tables) distributed across 3 storage tiers:


1. **database.db** — Main monolithic application database (**18 tables**, explicitly listed in [database.py:_MONOLITH_TABLES](#) lines 54-73). Contains all core tables: player statistics ([PlayerMatchStats](#)), coach state ([CoachState](#)), ingestion tasks ([IngestionTask](#)), coaching insights ([CoachingInsight](#)), user profiles ([PlayerProfile](#)), system notifications ([ServiceNotification](#)), RAG base ([TacticalKnowledge](#)), COPER experience bank


(`CoachingExperience`), match results (`MatchResult` , `MapVeto`), calibrations (`CalibrationSnapshot`), role thresholds (`RoleThresholdRecord`), provenance tracking (`DataLineage` , `DataQualityMetric`), and extended tables for team round (`Ext_TeamRoundStats`) and playstyle (`Ext_PlayerPlaystyle`), plus archival `RoundStats` and `PlayerTickState` .

2. `hltv_metadata.db` — Professional metadata database (**3 tables**, `database.py:_HLTV_TABLES` lines 78-82): pro player profiles (`ProPlayer` , `ProTeam`) and stat cards (`ProPlayerStatCard`). Kept separate from the monolith because it is written by a separate process (HLTV sync service) to eliminate WAL contention with the session engine daemons.
3. `match_data/{id}.db` — Per-match telemetry database (**3 tables**, defined in `match_data_manager.py` : `MatchTickState` :110, `MatchEventState` :190, `MatchMetadata` :242). Each match has its own dedicated SQLite file containing tick-by-tick data (~100,000 rows per match). Managed by `MatchDataManager` . This separation solves the "Telemetry Cliff" problem — it prevents the monolith database from growing indefinitely with high-frequency data.


This three-tier separation ensures that the session engine's intensive write operations (demo ingestion, ML training → `database.db`) do not contend for WAL locks with HLTV scraping in a separate process (`hltv_metadata.db`), and that high-frequency per-match telemetry does not bloat the monolith (`match_data/{id}.db`).







PlayerTickState
(~100,000 rows/match)
Position, health, weapon
every 1/128 second



ER diagram explanation: Each box represents a **record type** in the database. **PlayerMatchStats** is like a **report card** for each player in each match (how many kills, deaths, their rating, etc.).

PlayerTickState is like a **frame-by-frame diary**: 128 entries per second recording exactly where the player was, how healthy they were, what direction they were looking. **RoundStats** is the **per-question breakdown**: individual evaluations for each round (kills, deaths, damage, noscope kills, flash assists, round rating), enabling detailed analyses. **CoachingExperience** is the coach's **diary**: every coaching moment, whether the advice worked, and how effective it was. **CoachingInsight** is the **actual advice** given to the player. **TacticalKnowledge** is the **textbook**: tips and strategies the coach can consult. **RoleThresholdRecord** is the **grading rubric**, the learned thresholds for classifying player roles. **CalibrationSnapshot** is the **instrument control log**, recording when the belief model was recalibrated and with how many samples. **Ext_PlayerPlaystyle** is the **external scouting report**, the playstyle metrics derived from CSV data used to train NeuralRoleHead.

ServiceNotification is the **intercom system**, the error and event messages from background daemons shown in the UI. The lines between tables show relationships: every match record is linked to a player's profile, coaching experiences are linked to specific matches, and RoundStats is linked to PlayerMatchStats via demo_name. **DataLineage** is the **provenance log**: it traces which demo originated each entity and through which pipeline step it was processed (append-only, for full audit trail). **DataQualityMetric** is the **quality metrics panel**: it records numeric quality values for each pipeline run (e.g., percentage of samples discarded, zero-tensor fallback rate), enabling system health monitoring over time.

Data lifecycle:

Phase	Tables written	Volume
Demo ingestion	<code>PlayerMatchStats</code> , <code>PlayerTickState</code> , <code>MatchMetadata</code>	~100,000 ticks/match
Round enrichment	<code>RoundStats</code> (per-round, per-player isolation)	~30 rows/match (rounds × players)
HLTV scan	<code>ProPlayer</code> , <code>ProTeam</code> , <code>ProPlayerStatCard</code>	~500 players
CSV import	External tables via <code>csv_migrator.py</code>	~10,000 rows
Playstyle data	<code>Ext_PlayerPlaystyle</code> (from CSV for NeuralRoleHead)	~300+ players
Feature engineering	<code>PlayerMatchStats.dataset_split</code> updated	In place
RAG population	<code>TacticalKnowledge</code>	~200 articles
Experience extraction	<code>CoachingExperience</code>	~1,000 per match
Coaching output	<code>CoachingInsight</code>	~5-20 per match
Role threshold learning	<code>RoleThresholdRecord</code>	9 thresholds
Belief calibration	<code>CalibrationSnapshot</code> (after retraining)	1 per calibration run
System telemetry	<code>CoachState</code> , <code>ServiceNotification</code> , <code>IngestionTask</code>	Continuous
Provenance tracking	<code>DataLineage</code> (append-only per processed entity)	~N rows per ingested demo
Pipeline quality metrics	<code>DataQualityMetric</code> (append-only per pipeline run)	~5-10 metrics per run
Backup	Automated via <code>BackupManager</code> (7 daily rotations + 4 weekly)	Full database copy

Indexes and query optimization:

The most queried tables have strategic indexes to guarantee fast queries:

Table	Index	Columns	Query type optimized
PlayerMatchStats	idx_pms_player	player_name	Lookup by player
PlayerMatchStats	idx_pms_demo	demo_name	Lookup by match
PlayerMatchStats	idx_pms_processed	processed_at	Chronological sort
RoundStats	idx_rs_demo_round	demo_name, round_number	Lookup for specific round
IngestionTask	idx_it_status	status	Work queue (status=queued)
CoachingInsight	idx_ci_player	player_name	Insights per player
TacticalKnowledge	idx_tk_category	category	RAG lookup by category
ProPlayer	idx_pp_name	nickname	Pro lookup by name
DataLineage	idx_dl_entity	entity_type, entity_id	Provenance tracking per entity
DataQualityMetric	idx_dqm_run	run_id, run_type	Quality metrics per run

Integrity constraints:

Constraint	Tables	Enforcement
player_name NOT NULL	PlayerMatchStats, RoundStats	Schema-level
demo_name UNIQUE per player	PlayerMatchStats	Prevents duplicates
status CHECK IN ('queued', 'processing', 'completed', 'failed')	IngestionTask	Enum enforcement
dataset_split CHECK IN ('train', 'val', 'test')	PlayerMatchStats	Split validity
Foreign key demo_name	RoundStats → PlayerMatchStats	Round-match relationship

Key table details:

PlayerMatchStats (32 fields) — the most queried table in the system:

The PlayerMatchStats table contains all aggregated per-player, per-match statistics. It is the "report card" for every analyzed match:

Field	Type	Description
<code>id</code>	Integer PK	Unique identifier
<code>player_name</code>	String	Player name (indexed)
<code>demo_name</code>	String	Demo file name (unique per player)
<code>kills</code> , <code>deaths</code> , <code>assists</code>	Integer	Base KDA
<code>adr</code>	Float	Average Damage per Round
<code>kast</code>	Float	Kill/Assist/Survive/Trade %
<code>headshot_percentage</code>	Float	HS%
<code>hltv_rating</code>	Float	Computed HLTV 2.0 Rating
<code>dataset_split</code>	String	"train" / "val" / "test"
<code>is_pro</code>	Boolean	Pro player flag
<code>processed_at</code>	DateTime	Processing timestamp
<code>map_name</code>	String	Map played
...	...	20+ additional fields for advanced features

TacticalKnowledge — the RAG base:

Field	Type	Description
<code>id</code>	Integer PK	Identifier
<code>title</code>	String	Document title (indexed)
<code>description</code>	String	Description of tactical content
<code>category</code>	String	"positioning" / "economy" / "utility" / "aim" (indexed)
<code>map_name</code>	String	Specific map (indexed, optional)
<code>situation</code>	String	Situational context: "T-side pistol round", "CT retake A site"
<code>pro_example</code>	String	Reference to pro demo (optional)
<code>embedding</code>	String	384-dim JSON vector (sentence-transformers)
<code>created_at</code>	DateTime	Creation timestamp
<code>usage_count</code>	Integer	Usage counter

RAG semantic search works by computing the **cosine similarity** between the user query embedding and the precomputed embeddings of each document. The top-3 results with similarity > 0.5 are used to enrich the coaching context. The `situation` field enables contextual filtering (e.g., "T-side pistol round") before semantic search.

CoachingExperience — the COPER bank (22+ fields):


Field	Type	Description
<code>id</code>	Integer PK	Identifier
<code>context_hash</code>	String	Hash of game state for fast lookup (indexed)
<code>map_name</code>	String	Map (indexed)
<code>round_phase</code>	String	"pistol" / "eco" / "full_buy" / "force"
<code>side</code>	String	"T" / "CT"
<code>position_area</code>	String	"A-site" / "Mid" / etc. (indexed, optional)
<code>game_state_json</code>	String	Full tick snapshot (max 16KB, validated with <code>field_validator</code>)
<code>action_taken</code>	String	"pushed" / "held_angle" / "rotated" / "used_utility" / etc.
<code>outcome</code>	String	"kill" / "death" / "trade" / "objective" / "survived" (indexed)
<code>delta_win_prob</code>	Float	Win probability change from this action
<code>confidence</code>	Float	Reliability/generalizability 0.0-1.0
<code>usage_count</code>	Integer	How many times retrieved for coaching
<code>pro_match_id</code>	Integer FK	Reference to MatchResult (ON DELETE SET NULL)
<code>pro_player_name</code>	String	Reference pro player name (indexed, optional)
<code>embedding</code>	String	384-dim JSON vector for semantic search (optional)
<code>source_demo</code>	String	Source demo (optional)
<code>created_at</code>	DateTime	Creation timestamp
<code>outcome_validated</code>	Boolean	Whether the outcome was validated
<code>effectiveness_score</code>	Float	Score -1.0 to 1.0
<code>follow_up_match_id</code>	Integer	Follow-up match for tracking
<code>times_advice_given</code>	Integer	How many times the advice was given
<code>times_advice_followed</code>	Integer	How many times the advice was followed
<code>last_feedback_at</code>	DateTime	Last feedback received
<code>mu_skill</code>	Float	TrueSkill μ posterior — experience quality estimate (KT-01)
<code>sigma_skill</code>	Float	TrueSkill σ uncertainty — quality uncertainty (KT-01)
<code>times_retrieved</code>	Integer	Replay priority counter — how many times retrieved (KT-01)
<code>times_validated</code>	Integer	User confirmation counter — how many times validated (KT-01)

The COPER system uses these experiences to **learn from its own advice**: the `context_hash` field enables fast lookup of similar situations, `outcome` and `delta_win_prob` measure action effectiveness, and the feedback loop (`outcome_validated` , `times_advice_given/followed` , `effectiveness_score`) lets the system prioritize advice that has historically produced positive outcomes. The `game_state_json`


field is capped at 16KB to prevent uncontrolled database growth. The TrueSkill `mu_skill` / `sigma_skill` fields (KT-01) provide a Bayesian estimate of experience quality, used for replay prioritization:

$$\text{confidence_score} = \text{mu_skill} - \kappa \times \text{sigma_skill}$$

The `times_retrieved` counter implements replay priority, while `times_validated` tracks user confirmations for semantic CRUD.




10. Training regime and maturity gates




Data requirements per phase:

Phase	Minimum data	Training type	Primary loss
1. JEPA Pre-training	10 pro demos	Self-supervised (InfoNCE)	Contrastive with in-batch negatives
2. Pro baseline	50 pro matches	Supervised	MSE(pred, pro_stats)
3. User fine-tuning	50 user matches	Supervised (transfer)	MSE(pred, user_stats)
4. RAP optimization	200 total matches	Multi-task	Strategy + Value + Sparsity + Position



VL-JEPA Two-Stage Protocol (Concept Alignment):

When VL-JEPA is active, Phases 1-2 are extended with a **two-stage protocol** that aligns latent representations with the 16 coaching concepts:



Stage	What is trained	What is frozen	Loss	Purpose
1	Context encoder, predictor	Target encoder (EMA)	InfoNCE (τ=0.07)	General latent representations
2	Concept embeddings, concept projector, concept temperature	Encoder (optional fine-tuning)	BCE + VICReg diversity	Alignment with 16 coaching concepts

The 16 Coaching Concepts (taxonomy):


Index	Concept	Category	Description
0	Positioning Quality	Positioning	Position quality relative to context
1	Trade Readiness	Positioning	Readiness for trade kill
2	Rotation Speed	Positioning	Rotation speed between sites
3	Utility Usage	Utility	Grenade use frequency and quality
4	Utility Effectiveness	Utility	Effectiveness of utilities used
5	Decision Quality	Decision	In-game decision quality
6	Risk Assessment	Decision	Pre-action risk evaluation
7	Engagement Timing	Engagement	Timing of engagements
8	Engagement Distance	Engagement	Optimal engagement distance
9	Crosshair Placement	Engagement	Pre-peek crosshair placement
10	Recoil Control	Engagement	Recoil control
11	Economy Management	Decision	Team economy management
12	Information Gathering	Decision	Information gathering (peek, info utility)
13	Composure Under Pressure	Psychology	Composure under pressure
14	Aggression Control	Psychology	Aggression control
15	Adaptation Speed	Psychology	Speed of adaptation to opponent meta

AdamW + CosineAnnealing (JEPA Trainer):

Hyperparameter	Value	Purpose
Optimizer	AdamW	Weight decay separated from gradients
Learning rate	1e-4 (default)	Initial learning rate
Weight decay	0.01	L2 regularization
Scheduler	CosineAnnealingLR	Cosine LR decay down to 0
EMA decay	0.996	Target encoder momentum update
Gradient clip	1.0	Gradient explosion prevention

DriftMonitor (z_threshold=2.5):

The `DriftMonitor` integrated in the JEPA Trainer monitors **feature drift** during training. If a feature's Z-score exceeds 2.5, it emits a warning indicating possible distribution shift:



Retraining trigger: The Teacher daemon monitors growth in the pro demo count; it triggers retraining when `count ≥ last_count × 1.10`.

JEPAPretrainDataset (`jepa_train.py`):

The JEPA pre-training dataset uses **temporal windows** to create context-target pairs:

Parameter	Value	Purpose
<code>context_len</code>	10 ticks	Context window length (input)
<code>target_len</code>	10 ticks	Target window length (to predict)
Gap	0-5 ticks (random)	Variable distance between context and target
Batch size	32	Number of pairs per batch

11. Loss function catalog

Model	Loss name	Formula	Purpose
JEPA	InfoNCE Contrastive	$-\log(\exp(\text{sim}(\text{pred}, \text{target})/\tau) / \sum \exp(\text{sim}(\text{pred}, \text{neg}_i)/\tau))$, $\tau=0.07$, <code>F.normalize</code> before cosine similarity	Aligning context predictions with target embeddings
JEPA	Fine-tuning	<code>MSE(coaching_head(Z_ctx), y_true)</code>	Supervised coaching score
AdvancedCoachNN	Supervised	<code>MSELoss(MoE_output, y_true)</code>	Match-level training
RAP	Strategy	<code>MSELoss(advice_probs, target_strat)</code>	Correct tactical recommendation
RAP	Value	$0.5 \times \text{MSE}(V(s), \text{true_advantage})$	Accurate advantage estimation
RAP	Sparsity	<code>L1(gate_weights)</code>	Expert specialization
RAP	Position	<code>MSE(xy) + 2 \times \text{MSE}(z)</code>	Optimal positioning with Z-axis penalty
WinProb	Prediction	<code>BCEWithLogitsLoss(pred, outcome)</code>	Round outcome prediction
NeuralRoleHead	KL-Divergence	<code>KLDivLoss(log_softmax(pred), target)</code> with label smoothing $\epsilon=0.02$	Role probability distribution match
VL-JEPA	Concept Alignment	<code>BCE(concept_logits, concept_labels) + \text{VICReg}(\text{concept_diversity})</code>	Visual-language concept grounding

Detail: InfoNCE Contrastive Loss (JEPA)

InfoNCE is the main loss for JEPA pre-training. Its purpose is to align context predictions with target embeddings while simultaneously pushing away negatives (other samples in the batch):

$$L_{\text{InfoNCE}} = -\log\left(\frac{\exp(\text{sim}(\text{pred}, \text{target}^+))}{\tau} / \sum_i \frac{\exp(\text{sim}(\text{pred}, \text{target}_i))}{\tau}\right)$$

Component	Value	Role
<code>sim()</code>	Cosine similarity after <code>F.normalize</code>	Similarity measure [-1, +1]
τ (temperature)	0.07	Distribution sharpness — low values = more selective
<code>target⁺</code>	The correct target embedding for this context	The "positive" — the correct answer
<code>target_i</code>	All embeddings in the batch	In-batch negatives — the wrong answers
Batch size	32	Number of negatives = <code>batch_size - 1</code> = 31

Detail: VL-JEPA Concept Loss (VICReg Components)

The VL-JEPA concept alignment loss combines two components:


```
L_concept = BCE(concept_logits, concept_labels) + α·VICReg_diversity
```

Where `VICReg_diversity` is composed of:

VICReg Term	Formula	Weight	Purpose
Variance	<code>max(0, γ - std(z))</code> for each dimension	$\alpha=0.5$	Prevents representation collapse — each dimension must vary
Covariance	<code>$\sum_{i \neq j} \text{cov}(z_i, z_j)^2$</code>	$\beta=0.1$	Decorrelation — each dimension must capture different information

Detail: RAP Multi-Task Loss


The RAP Coach combines 4 losses into a total weighted loss:



The 2× Z penalty in position loss reflects the fact that in CS2 getting the **floor** wrong (upper vs lower in Nuke/Vertigo) is much worse than getting horizontal position wrong. A 100-unit error on the Z axis (wrong floor) is strategically catastrophic, while the same error on X/Y could be irrelevant.

12. Complete Program Logic — From Launch to Advice

This chapter documents the **complete logic** of Macena CS2 Analyzer, from the moment the user launches the application to when they receive coaching advice. Unlike the previous chapters that focus on AI subsystems, here it is explained how **every component of the program** works together: the desktop interface, the quad-daemon architecture, the ingestion pipeline, the storage system, tactical playback, observability, and the application lifecycle.



12.1 Entry Point and Startup Sequence

The system has **two main entry points** (Qt primary, Kivy legacy) and three utility entry points:

#	Entry Point	Command	Role
1	Qt (primary)	<code>python -m Programma_CS2_RENAN.apps.qt_app.app</code>	PySide6 desktop UI
2	Kivy (legacy)	<code>python -m Programma_CS2_RENAN.main</code>	Kivy/KivyMD desktop UI
3	Session Engine	<code>python -m Programma_CS2_RENAN.backend.console</code>	Headless backend
4	Headless Validator	<code>python -m Programma_CS2_RENAN.tools.headless_validator</code>	CI/CD validation
5	HLTV Sync	<code>python -m Programma_CS2_RENAN.tools.hltv_sync</code>	Pro data scraping

12.1.1 Qt Entry Point (Primary) — apps/qt_app/app.py

File: `Programma_CS2_RENAN/apps/qt_app/app.py`

The Qt startup sequence follows a more modern approach based on **QApplication** and **signal/slot** instead of the Kivy event loop:

- High-DPI setup** — Enables automatic scaling for high-density displays
- QApplication** — Creates the Qt application instance with args handling
- Theme and fonts** — Registers the 3 themes (CS2, CSGO, CS1.6) via QSS and the custom fonts
- MainWindow** — Builds `QMainWindow` with sidebar + `QStackedWidget` (14 screens)
- Signal wiring** — Connects Qt signals between sidebar, screens, and backend
- First-run gate** — If `SETUP_COMPLETED=False`, shows the wizard; otherwise the home
- Backend console boot** — Launches the Session Engine as a subprocess
- Window show** — Shows the window and starts the Qt event loop


9. **CoachState polling** — Qt timer for periodic state updates


12.1.2 Kivy Entry Point (Legacy) — `main.py`

File: `Programma_CS2_RENAN/main.py`

When the user launches the application via the legacy entry point, `main.py` orchestrates a strictly ordered **9-phase startup sequence**. Each phase must complete successfully before the next can begin. If a critical phase fails, the application terminates with an explicit message — never silently.





**Critical phase details:**

Phase	Component	What it does	Failure consequence
1	<code>integrity.py</code> (RASP)	Verifies source file hashes against manifest	Immediate termination — possible tampering
2	<code>config.py</code>	Stabilizes <code>sys.path</code> , defines all path constants	Cascading import errors
3	<code>db_migrate.py</code>	Runs pending Alembic migrations	Incompatible schema — crash on DB ops
5	Kivy <code>Config</code>	Sets <code>KIVY_NO_CONSOLELOG</code> , registers fonts, loads <code>.kv</code>	UI not renderable
6	<code>database.py</code>	<code>create_all()</code> on SQLite engine with <code>check_same_thread=False</code>	No persistence possible
7	<code>lifecycle.py</code>	Launches subprocess with correct PYTHONPATH, checks mutex	No background automation


12.2 Lifecycle Management (`lifecycle.py`)

File: `Programma_CS2_RENAN/core/lifecycle.py`

The `AppLifecycleManager` is a **Singleton** that manages the lifecycle of the entire application: from guaranteeing that only one instance is active, to launching the daemon subprocess, to coordinated shutdown.

Main mechanisms:

Mechanism	Implementation	Purpose
Single Instance Lock	Windows Named Mutex / file lock on Linux	Prevents multiple instances (DB corruption)
Daemon Launch	<code>subprocess.Popen(session_engine.py)</code> with <code>PYTHONPATH</code>	Separate process for heavy work
Parent Death Detection	The daemon monitors EOF on <code>stdin</code>	If the main process dies, the daemon shuts down
Graceful Shutdown	Sends "STOP" via <code>stdin</code> → daemon terminates within 5s	No data loss or zombie tasks
Status Polling	UI queries <code>CoachState</code> every 10s	Daemon state update without direct IPC
Error Recovery	If daemon dies, error logged + ServiceNotification	User informed, no silent crash
Keep-Alive	UI checks heartbeat every 15s	If heartbeat > 30s stale → warning "Daemon not responding"





12.3 Configuration System (`config.py`)

File: `Programma_CS2_RENAN/core/config.py`

The system uses **three configuration levels**, each with a different level of persistence and security:







Critical system constants:

Constant	Value	File	Purpose
METADATA_DIM	25	config.py	Feature vector dimension — unified contract
SCAN_INTERVAL	10	config.py	Scanner scan interval (seconds)
MAX_DEMOS_PER_MONTH	10	config.py	Monthly demo upload quota
MAX_TOTAL_DEMOS	100	config.py	Lifetime total demo limit
MIN_DEMOS_FOR_COACHING	10	config.py	Threshold for full personalized coaching
TRADE_WINDOW_TICKS	192	trade_kill_detector.py	Trade kill time window (~3 seconds at 64 tick)
HLTV_BASELINE_KPR	0.679	demo_parser.py	HLTV 2.0 baseline for KPR
HLTV_BASELINE_SURVIVAL	0.317	demo_parser.py	HLTV 2.0 baseline for survival
FOV_DEGREES	90	player_knowledge.py	Simulated player field of view
MEMORY_DECAY_TAU	160	player_knowledge.py	Memory decay constant in ticks
CONFIDENCE_ROUNDS_CEILING	300	correction_engine.py	Round ceiling for max confidence
SILENCE_THRESHOLD	0.2	explainability.py	Threshold below which silence is valid action
MIN_SAMPLES_FOR_VALIDITY	10	role_thresholds.py	Minimum samples for a valid role threshold
HALF_LIFE_DAYS	90	pro_baseline.py	Temporal decay for pro data
Z_LEVEL_THRESHOLD	200	connect_map_context.py	Z threshold for floor classification

Path architecture:


The system handles paths with particular attention to **Windows/Linux portability**. The heart is `BRAIN_DATA_ROOT`: a user-configurable directory that contains models, logs, and derived data. If it doesn't exist, the system falls back to the project folder.

Path	Contents	Configurable
<code>DATABASE_URL</code>	Main monolithic database (<code>database.db</code>)	No — always in the project folder
<code>BRAIN_DATA_ROOT</code>	Root for derived data (models, logs)	Yes — via <code>user_settings.json</code>
<code>MODELS_DIR</code>	Model checkpoints <code>.pt</code>	Derived from <code>BRAIN_DATA_ROOT</code>
<code>LOG_DIR</code>	Application log files	Derived from <code>BRAIN_DATA_ROOT</code>
<code>MATCH_DATA_PATH</code>	Per-match database (<code>match_XXXX.db</code>)	Derived from <code>BRAIN_DATA_ROOT</code>
<code>DEFAULT_DEMO_PATH</code>	User demo folder	Yes — via UI Settings
<code>PRO_DEMO_PATH</code>	Pro demo folder	Yes — via UI Settings

12.4 Session Engine — Quad-Daemon Architecture (`session_engine.py`)

File: `Programma_CS2_RENAN/core/session_engine.py`


The Session Engine is the **beating heart** of the system's automation. It lives as a separate subprocess and hosts **4 daemon threads** that work in parallel, each with a well-defined responsibility. This design completely separates the heavy work (demo parsing, ML training) from the user interface, guaranteeing that the Kivy GUI always remains responsive.



Lifecycle of each daemon:

Daemon	Interval	Work per cycle	Trigger
Scanner	10 seconds	Scans pro and user folders, creates <code>IngestionTask</code> for new <code>.dem</code>	Always active (if state = Scanning)
Digester	Continuous	Picks 1 task from queue, performs full parsing	<code>_work_available_event</code> (signaled by Scanner)
Teacher	300 seconds (5 min)	Checks pro sample growth; if $\geq 10\%$ → <code>run_full_cycle()</code>	<code>pro_count >= last_count × 1.10</code>
Pulse	5 seconds	Updates <code>CoachState.last_heartbeat</code> in the database	Always active

Shutdown sequence:



Daemon implementation details:

Daemon	Main Method	Internal Loop	Exit Condition
Scanner	<code>_scanner_daemon_loop()</code>	<code>while not _shutdown_event.is_set() → scan_all_paths() → sleep(10)</code>	<code>_shutdown_event</code> set
Digester	<code>_digester_daemon_loop()</code>	<code>while not _shutdown_event.is_set() → _work_available_event.wait(2) → process_next_task()</code>	<code>_shutdown_event</code> set
Teacher	<code>_teacher_daemon_loop()</code>	<code>while not _shutdown_event.is_set() → _check_retrain_needed() → sleep(300)</code>	<code>_shutdown_event</code> set
Pulse	<code>_pulse_daemon_loop()</code>	<code>while not _shutdown_event.is_set() → state_mgr.heartbeat() → sleep(5)</code>	<code>_shutdown_event</code> set

Error handling in daemons:

Each daemon is protected by a global `try/except`. If a daemon crashes:

1. The error is logged with a full traceback
2. The `StateManager` registers the error (`set_error(daemon, message)`)
3. A `ServiceNotification` is created for the user
4. The daemon is **not automatically restarted** (by design: a daemon crash indicates a bug, not a transient error)
5. The other daemons continue to function independently

Zombie Task Cleanup: On startup, the Session Engine looks for tasks with `status="processing"` left over from a previous crash and resets them to `status="queued"`, enabling automatic recovery without data loss.

Automatic Backup: On Session Engine startup, `BackupManager.should_run_auto_backup()` checks whether a backup is needed and, if so, creates a checkpoint labeled `"startup_auto"`. Backups follow a rotation of 7 daily + 4 weekly copies.

12.5 Desktop Interface

The desktop interface has two implementations: **Qt/PySide6** (primary) and **Kivy/KivyMD** (legacy). Both follow the **MVVM** pattern (Model-View-ViewModel).

12.5.1 Qt Interface (Primary) — `apps/qt_app/`

Directory: `Programma_CS2_RENAN/apps/qt_app/` **Key files:** `app.py`, `main_window.py`, `core/i18n_bridge.py`, `core/theme_engine.py`, `screens/`

The Qt interface is built with **PySide6 (Qt 6)** and uses an **MVVM pattern with Qt Signals/Slots**. The `MainWindow` (`QMainWindow`) is composed of a **navigation sidebar** and a `QStackedWidget` that hosts the 14 screens.

Specification	Detail
Framework	PySide6 (Qt 6 for Python)
Pattern	MVVM with Qt Signals/Slots
Platforms	Windows, macOS, Linux
Resolution	Adaptive, native High-DPI
Themes	3: CS2 (orange), CSGO (blue-gray), CS1.6 (green) — QSS + QPalette
i18n	3 languages: EN, IT, PT — JSON + <code>QtLocalizationManager</code>
Charts	QPainter for tactical map, native Qt widgets for charts

i18n system: The `QtLocalizationManager` (`core/i18n_bridge.py`) loads per-language JSON files (`en.json`, `it.json`, `pt.json`) and handles runtime language switching via Qt signals. Every UI string is dynamically resolved via a localization key.


Theme system: The `ThemeEngine` (`core/theme_engine.py`) applies QSS stylesheets and configures Qt's `QPalette` for each theme:

- **CS2** — Orange palette (#FF6600) with dark background, inspired by the CS2 UI
- **CSGO** — Blue-gray palette (#4A90D9) with cool tones, inspired by CS:GO
- **CS1.6** — Green palette (#33CC33) on black background, inspired by the retro look of CS 1.6

12.5.2 Kivy Interface (Legacy) — `apps/desktop_app/`

Directory: `Programma_CS2_RENAN/apps/desktop_app/` **Key files:** `layout.kv`, `wizard_screen.py`, `player_sidebar.py`, `tactical_viewer_screen.py`, `tactical_viewmodels.py`, `tactical_map.py`, `timeline.py`, `widgets.py`, `help_screen.py`, `ghost_pixel.py`

The legacy interface is built with **Kivy + KivyMD** and follows the **MVVM** pattern (Model-View-ViewModel). The `ScreenManager` handles navigation between screens with `FadeTransition` transitions.



14 interface screens:

Screen	Role	Key components
Wizard	First-time setup	Player name, role, demo folder paths
Home	Dashboard	Monthly quota (X/10), service status (green/red), belief confidence (0-1), active tasks, processed matches counter
Coach	Coaching insights	Color-coded cards by severity, multi-dimensional skill radar, historical trends, AI chat (Ollama/Claude), active tasks
Tactical Viewer	Tactical playback	2D map with players/grenades/ghost, timeline with event markers, CT/T player sidebar, speed controls (0.25x → 8x)
Settings	Customization	Theme (CS2/CSGO/CS1.6), font, text size, language, demo paths, wallpaper
Help	User support	Interactive tutorial, FAQ, troubleshooting
Match History	Match archive	List of analyzed demos with filters and sorting
Match Detail	Match detail	Detailed statistics for a single analyzed demo
Performance	Progress	5-axis skill radar, trend charts, temporal comparisons
User Profile	User profile	Bio, preferred role, Steam/FACEIT sync
Profile	Public profile	Public player profile view
Steam Config	Steam configuration	Entering and validating Steam API key
Pro Comparison	Pro comparison	User statistics comparison with HLTV pro players, performance benchmark
FACEIT Config	FACEIT configuration	Entering and validating FACEIT API key

Custom widgets:

Widget	File	Function
PlayerSidebar	player_sidebar.py	CT/T list with role icons, health/armor, current weapon, money, and alive/dead status
TacticalMap	tactical_map.py	2D canvas with multi-layer rendering: map texture → heatmap → players → grenades → ghost (QPainter in Qt, Canvas Kivy in legacy)
Timeline	timeline.py	Horizontal scrubber with tick numbers, color-coded event markers, drag-to-see, double-click jump
GhostPixel	ghost_pixel.py	Rendering of the semi-transparent ghost circle (optimal position predicted by RAP)


Available themes:


The application supports **3 themes** selectable from the Settings screen, each with a custom color palette and wallpaper:

Theme	Primary palette	Inspiration
CS2 (default)	Steel blue + orange	Modern Counter-Strike 2 UI
CSGO	Military green + yellow	Counter-Strike: Global Offensive
CS 1.6	Dark brown + lime green	Classic Counter-Strike 1.6 (nostalgia)

Coach Screen — Detailed layout:


The Coach screen is the most complex in the application, with 5 functional areas:



MVVM pattern in the Tactical Viewer:**12.6 Ingestion Pipeline (`ingestion/`)**

Directory: `Programma_CS2_RENAN/ingestion/` **Key files:** `demo_loader.py` , `steam_locator.py` , `integrity.py` , `registry/` , `pipelines/user_ingest.py` , `pipelines/json_tournament_ingestor.py`

The ingestion pipeline is the **complete path** a `.dem` file takes from the filesystem to becoming coaching insights in the database. It is orchestrated by the Scanner daemon (discovery) and the Digester daemon (processing).



DemoLoader — The parser at the heart of the pipeline:

The `DemoLoader` is the wrapper around `demoparser2` (high-performance Rust library) that transforms a binary `.dem` file into Python data structures:

Parsing Phase	Output	Typical size
1. Header parsing	Metadata (map, server, duration)	~100 bytes
2. Frame extraction	List of frames (position, health, weapon per tick)	~100,000 frames/match
3. Event extraction	List of events (kill, death, bomb_plant, round_start, etc.)	~500-2,000 events/match
4. Player summary	Aggregated per-player statistics	~10 records

FeatureExtractor (`backend/processing/feature_engineering/vectorizer.py`):

The FeatureExtractor is the component that transforms raw demo data into **25-dimensional numeric vectors** (`METADATA_DIM=25`) usable by the neural networks. **Important:** these are **tick-level** features (128 Hz), not match-level aggregated statistics. Each individual game frame produces a 25-dim vector that captures the instantaneous state of the player:

Dim	Feature	Type	Range	Description
0	health	Float	[0, 1]	Normalized health
1	armor	Float	[0, 1]	Normalized armor
2	has_helmet	Binary	0/1	Helmet equipped
3	has_defuser	Binary	0/1	Defuse kit equipped
4	equipment_value	Float	[0, 1]	Normalized equipment value
5	is_crouching	Binary	0/1	Crouched
6	is_scoped	Binary	0/1	Scope active
7	is_blinded	Binary	0/1	Flashed
8	enemies_visible	Float	[0, 1]	Visible enemies (normalized, clamped)
9	pos_x	Float	[-1, 1]	X position (normalized $\pm$ pos_xy_extent)
10	pos_y	Float	[-1, 1]	Y position (normalized $\pm$ pos_xy_extent)
11	pos_z	Float	[0, 1]	Z position (normalized, handles Nuke/Vertigo)
12	view_x_sin	Float	[-1, 1]	sin(yaw) — cyclic continuity of horizontal angle
13	view_x_cos	Float	[-1, 1]	cos(yaw) — cyclic continuity of horizontal angle
14	view_y	Float	[-1, 1]	Normalized pitch (vertical angle)
15	z_penalty	Float	[0, 1]	Vertical level distinction (floor penalty)
16	kast_estimate	Float	[0, 1]	KAST estimate (participation ratio)
17	map_id	Float	[0, 1]	Deterministic map hash
18	round_phase	Float	{0, 0.33, 0.66, 1}	Economic phase: pistol/eco/force/full_buy
19	weapon_class	Float	{0–1.0}	Weapon class: 0=knife, 0.2=pistol, 0.4=SMG, 0.6=rifle, 0.8=sniper, 1.0=heavy
20	time_in_round	Float	[0, 1]	Seconds in round / 115 (clamped)
21	bomb_planted	Binary	0/1	Bomb planted
22	teammates_alive	Float	[0, 1]	Teammates alive (count / 4)
23	enemies_alive	Float	[0, 1]	Enemies alive (count / 5)
24	team_economy	Float	[0, 1]	Team money average / 16000 (clamped)


Normalization and bounds:

Normalization is integrated directly into the FeatureExtractor, with bounds configurable via `HeuristicConfig` (externalized to JSON). The cyclic encoding of view angles (sin/cos for yaw) prevents discontinuities at the 0°/360° boundary. The `z_penalty` automatically distinguishes floors on multi-level

maps (Nuke, Vertigo). Weapon class uses an ordinal categorical mapping (6 classes) defined in the `WEAPON_CLASS_MAP` constant (also includes grenades=0.1 and special equipment=0.05).

Temporal Dataset Split:

The dataset split follows a strict **chronological ordering** to prevent temporal data leakage:



Enrich From Demo — Post-parsing enrichment:

After base parsing, `enrich_from_demo()` adds advanced metrics computed from events:

Enriched Metric	Computation	Source
Trade kills	<code>TradeKillDetector.detect()</code> with <code>TRADE_WINDOW_TICKS=192</code>	Kill/death events
Flash assists	Count of blinds within time window before a kill	Blind + kill events
Noscope kills	Kills with sniper weapon without scope active	Kill event + weapon state
Wallbang kills	Kills through penetrable surfaces	Kill event with penetration flag
Through-smoke kills	Kills with smoke active in line of fire	Kill event + smoke position
Blind kills	Kills while the player is flashed	Kill event + flash state

Specific components:

Component	File	Role
DemoLoader	<code>demo_loader.py</code>	Wraps <code>demoparser2</code> , extracts frames and events from the <code>.dem</code> file
SteamLocator	<code>steam_locator.py</code>	Automatically locates the CS2 demo folder via Steam registry / <code>libraryfolders.vdf</code>
IntegrityChecker	<code>integrity.py</code>	Verifies that demo files are valid, complete, and not corrupt before parsing
UserIngestPipeline	<code>pipelines/user_ingest.py</code>	Complete pipeline for user demos: parse → enrich → stats → coaching
JsonTournamentIngestor	<code>pipelines/json_tournament_ingestor.py</code>	Imports tournament data from structured JSON files
Registry	<code>registry/registry.py</code>	Tracks all processed demos, prevents duplicates
ResourceManager	<code>ingestion/resource_manager.py</code>	Hardware resource management: CPU/RAM throttling, disk space
JsonTournamentIngestor	<code>pipelines/json_tournament_ingestor.py</code>	Imports tournament data from structured JSON files
RegistryLifecycle	<code>registry/lifecycle.py</code>	Lifecycle management of ingestion records

SteamLocator (`ingestion/steam_locator.py`) — automatic CS2 demo location:

The SteamLocator implements a **cross-platform discovery** algorithm to automatically find the CS2 demo folder:

Platform	Strategy	Typical path
Windows	System registry → <code>libraryfolders.vdf</code>	<code>C:\Program Files (x86)\Steam\steamapps\common\Counter-Strike Global Offensive\game\csgo\replays</code>
Linux	<code>~/.steam/steam/</code> → <code>libraryfolders.vdf</code>	<code>~/.steam/steam/steamapps/common/...</code>
Fallback	Prompts user via UI Settings	Custom path

IntegrityChecker (`ingestion/integrity.py`):


Preliminary check of each demo file before expensive parsing:

- **Magic bytes:** `PBDEMS2\0` (CS2 Source 2) or `HL2DEMO\0` (legacy Source 1)
- **Size bounds:** minimum 1KB (not empty), maximum 5GB (not corrupt/excessive)
- **Read test:** attempts to read the first N bytes to verify the file is accessible

12.7 Unified Control Console (`backend/control/`)

File: `Programma_CS2_RENAN/backend/control/console.py` , `ingest_manager.py` , `db_governor.py` , `ml_controller.py`

The Console is a **Singleton** that serves as a central coordination point for all backend subsystems. It is the "control panel" through which every part of the system can be controlled.




Console boot sequence:

1. `ServiceSupervisor` starts the "hunter" service (HLTV sync) as a monitored process
2. `DatabaseGovernor` performs an integrity audit: verifies all tables, counts records, checks schema
3. `MLController` remains on standby — training is handled by the Teacher daemon
4. `IngestionManager` stays idle — active work is handled by the Scanner/Digester daemons

MLControlContext — Live Training Control:

The `MLController` uses a control token called `MLControlContext` (`backend/control/ml_controller.py`) that is passed to training cycles to enable **real-time intervention** by the operator. This replaces the previous approach based on `StopIteration` with a more robust thread-safe system.




Command	Method	Effect
Pause	<code>request_pause()</code>	<code>_resume_event.clear()</code> → blocks <code>check_state()</code>
Resume	<code>request_resume()</code>	<code>_resume_event.set()</code> → unblocks training
Stop	<code>request_stop()</code>	Raises <code>TrainingStopRequested</code> (custom exception)
Throttle	<code>set_throttle(factor)</code>	Adds <code>time.sleep(factor)</code> after each batch

12.8 Onboarding and New User Flow

File: `Programma_CS2_RENAN/backend/onboarding/new_user_flow.py`

The `OnboardingManager` guides new users through a **3-phase progression** that automatically adapts to the amount of available data.



Wizard Screen (First-time Setup):

On first run (`SETUP_COMPLETED = False`), the user is guided through the Wizard:

1. **Player name** — The name that will appear in analyses
2. **Preferred role** — Entry Fragger, AWPer, Lurker, Support, or IGL
3. **Demo folder paths** — Where the system automatically looks for demos
4. On completion: `SETUP_COMPLETED = True` , redirect to Home

Quota cache (Task 2.16.1): The demo count is cached for 60 seconds to avoid repeated DB queries.

`invalidate_cache()` is called after every new upload, ensuring the UI always shows the correct count without overloading the database.

Help System (`backend/knowledge/help_system.py`):

The integrated help system provides contextual support to the user:


Feature	Implementation
Interactive tutorial	Step-by-step guides for the main features
Contextual FAQ	Frequently asked questions filtered by current screen
Troubleshooting	Decision tree for common problems (Steam path not found, demo not parsed, empty coaching)
Tooltips	Inline explanations for complex metrics (KAST, HLTV 2.0 Rating, ADR)


12.9 Storage Architecture (`backend/storage/`)

Directory: `Programma_CS2_RENAN/backend/storage/` **Key files:** `database.py` , `db_models.py` , `match_data_manager.py` , `storage_manager.py` , `maintenance.py` , `state_manager.py` , `stat_aggregator.py` , `backup.py`

The storage system uses a **three-tier storage** architecture based on SQLite in WAL mode (Write-Ahead Logging), which enables concurrent reads and writes without blocking.







The 21 SQLModel tables:

#	Table	Database	Category	Description
1	PlayerMatchStats	database.db	Core	Aggregated per-player/match statistics (32 fields)
2	PlayerTickState	database.db	Core	Per-tick state (128 Hz), also archived in separate per-match DBs
3	PlayerProfile	database.db	User	User profile (name, role, Steam ID, monthly quota)
4	RoundStats	database.db	Core	Per-round isolated statistics (kills, rating, enrichment)
5	CoachingInsight	database.db	Coaching	Advice generated by the coaching service
6	CoachingExperience	database.db	Coaching	COPER experience bank (context, outcome, effectiveness, TrueSkill μ/σ , replay priority — KT-01)
7	IngestionTask	database.db	System	Work queue for the Digester daemon
8	CoachState	database.db	System	Global state (training metrics, heartbeat, status)
9	ServiceNotification	database.db	System	Error/event messages from daemons → UI
10	TacticalKnowledge	database.db	Knowledge	RAG base (384-dim embedding in JSON)
11	ProPlayer	hlvtv_metadata.db	Pro	Pro player profiles
12	ProTeam	hlvtv_metadata.db	Pro	Pro team metadata
13	ProPlayerStatCard	hlvtv_metadata.db	Pro	Seasonal statistics per pro player
14	Ext_PlayerPlaystyle	database.db	External	Playstyle data from CSV (for NeuralRoleHead)
15	Ext_TeamRoundStats	database.db	External	External tournament statistics
16	MatchResult	database.db	Matches	Match outcomes
17	MapVeto	database.db	Matches	Map selection history
18	CalibrationSnapshot	database.db	System	Belief model calibration log (timestamp, samples, outcome)
19	RoleThresholdRecord	database.db	System	Learned thresholds for role classification (persistent across restarts)
20	DataLineage	database.db	Provenance	Append-only data provenance log: entity_type, entity_id, source_demo, pipeline_version, processing_step
21	DataQualityMetric	database.db	Provenance	Append-only quality metrics per run: run_id, run_type, metric_name, metric_value, sample_count

Support enums (not tables):

Enum	Type	Description
<code>DatasetSplit</code>	<code>str</code> , Enum	Split categories (train/val/test/unassigned) — used as constraint on <code>PlayerMatchStats.dataset_split</code>
<code>CoachStatus</code>	<code>str</code> , Enum	Coach states (Paused/Training/Idle/Error) — used as constraint on <code>CoachState.status</code>

Detailed Storage Components:

MatchDataManager (`backend/storage/match_data_manager.py`) — the largest component of the storage layer:

The MatchDataManager is responsible for managing high-density per-match data (PlayerTickState with ~100,000 rows per match). To prevent the main database from growing uncontrollably, each match has its own separate SQLite database (`match_XXXX.db`).

Method	Description
<code>create_match_db(demo_name)</code>	Creates a new per-match database with <code>PlayerTickState</code> schema
<code>store_tick_data(demo_name, ticks)</code>	Bulk insert of tick data into the dedicated DB
<code>load_match_frames(demo_name)</code>	Loads all frames for the Tactical Viewer
<code>get_match_db_path(demo_name)</code>	Resolves the path of the per-match DB
<code>list_available_matches()</code>	Lists all matches with DB available
<code>delete_match_data(demo_name)</code>	Removes the per-match DB and updates the registry
<code>get_match_statistics(demo_name)</code>	Computes aggregate statistics from tick data

StorageManager (`backend/storage/storage_manager.py`):

The StorageManager is the **high-level storage coordinator** that manages quotas, uploads, and data lifecycle:

Responsibility	Implementation
Monthly quota	<code>can_user_upload()</code> → checks <code>MAX_DEMOS_PER_MONTH=10</code> and <code>MAX_TOTAL_DEMOS=100</code>
Upload flow	<code>handle_demo_upload(path)</code> → validation → copy to working dir → create IngestionTask
Cleanup	<code>cleanup_old_data(days)</code> → removes old match DBs and tasks
Disk space	<code>get_storage_usage()</code> → reports size for each database and directory

StateManager (`backend/storage/state_manager.py`):

The StateManager is a **Singleton** that maintains the system's runtime state and persists it to the database via the `CoachState` table:

Method	Purpose
<code>update_status(daemon, text)</code>	Updates the state of a specific daemon
<code>heartbeat()</code>	Updates the <code>last_heartbeat</code> timestamp
<code>get_state()</code>	Returns the current state as a <code>CoachState</code> object
<code>set_error(daemon, message)</code>	Registers an error for a daemon with timestamp
<code>update_training_metrics(epoch, loss, val_loss, eta)</code>	Updates training metrics in real time
<code>get_belief_confidence()</code>	Returns the belief model's confidence level (0.0-1.0)

StatAggregator (`backend/storage/stat_aggregator.py`):

Computes aggregate statistics from raw per-round data:

Aggregation	Formula	Use
<code>avg_kills</code>	<code>mean(RoundStats.kills)</code>	Dashboard, radar chart
<code>avg_adr</code>	<code>mean(RoundStats.damage_dealt / rounds)</code>	Pro comparison
<code>avg_kast</code>	<code>mean(rounds_with_kast / total_rounds)</code>	HLTV metric
<code>accuracy</code>	<code>sum(hits) / sum(shots_fired)</code>	Mechanical performance
<code>trade_kill_rate</code>	<code>trade_kills / team_deaths</code>	Team play

BackupManager (`backend/storage/backup.py`):


Feature	Detail
Daily rotation	7 copies — the oldest is overwritten
Weekly rotation	4 copies — additional weekly backup
Automatic trigger	On Session Engine startup via <code>should_run_auto_backup()</code>
Manual trigger	Via Console or UI Settings
Format	Full copy of the <code>.db</code> file (not SQL dump)
Labeling	<code>startup_auto</code> , <code>manual</code> , <code>pre_migration</code>

Maintenance (`backend/storage/maintenance.py`):

Operation	Frequency	Purpose
<code>vacuum()</code>	Monthly	Compacts the database, reclaims space from deleted records
<code>analyze()</code>	After every bulk insert	Updates SQLite query optimizer statistics
<code>wal_checkpoint()</code>	On startup	Forces merge of WAL into the main database
<code>integrity_check()</code>	On startup	<code>PRAGMA integrity_check</code> — verifies structural consistency

DbMigrate (`backend/storage/db_migrate.py`):

Wrapper around Alembic that automates the execution of migrations:




Connection Pooling and Concurrency:

Parameter	Value	Purpose
<code>check_same_thread</code>	<code>False</code>	Enables multi-thread access
<code>timeout</code>	30 seconds	Busy timeout for WAL contention
<code>pool_size</code>	1	Single SQLite writer (single-writer safety)
<code>max_overflow</code>	4	Overflow connections for load spikes
WAL mode	Enabled	Unlimited concurrent reads

12.10 Playback Engine and Tactical Viewer

File: `Programma_CS2_RENAN/core/playback.py` , `playback_engine.py` ,
`apps/desktop_app/tactical_viewer_screen.py` , `tactical_map.py` , `timeline.py` ,
`player_sidebar.py`

The tactical playback system allows the user to **relive their matches** on an interactive 2D map, with AI overlay (optimal ghost position), event markers (kills, bomb plants), and full playback controls.




PlaybackEngine — Internal architecture:

The PlaybackEngine handles frame-by-frame playback with temporal interpolation. In the Qt implementation, the update timer uses `QTimer` instead of `Kivy Clock.schedule_interval`, keeping the same interpolation and buffering logic:

Feature	Detail
Frame rate	60 FPS (interpolation from demo's native 64-tick)
Variable speed	0.25x (slow-mo), 0.5x, 1x (normal), 2x, 4x, 8x (fast-forward)
Interpolation	Linear between adjacent frames for smooth movement
Buffering	Pre-loads 120 frames ahead to avoid lag
Seek	Direct access to any tick via index
Round selection	Jumps directly to the start of a specific round

GhostEngine — Real-time inference:

The GhostEngine is the component that transforms the RAP Coach's predictions into **ghost positions visible on the map**:



ChronovisorScanner — Critical moment detection:

The ChronovisorScanner analyzes the entire match and identifies the **decisive moments** based on advantage change:

Moment type	Condition	Marker color
Critical error	Death while in numerical advantage (e.g., 4v3 → 3v3)	Red
Excellent play	Kill while in numerical disadvantage (e.g., 2v3 → 2v2)	Green
Bomb plant	bomb_planted event with timer	Yellow
Clutch	1vN victory ($N \geq 2$)	Gold
Eco round win	Victory with equipment < \$2,000	Blue

Each moment is placed on the Timeline as a clickable marker. Clicking jumps playback directly to that tick.


Viewer loading flow (One-Click):

1. The user clicks "Tactical Viewer" from Home
2. If no demo is loaded → `trigger_viewer_picker()` automatically opens the file picker
3. The user selects a `.dem` file
4. A "Dynamic 2D Reconstruction in Progress..." dialog appears
5. A background thread runs `_execute_viewer_parse(path)` → `DemoLoader.load_demo()`
6. On completion: dialog dismissal, loading frames into PlaybackEngine
7. The map is rendered with players at frame 0

12.11 Spatial Data and Map Management

File: `Programma_CS2_RENAN/core/spatial_data.py`, `spatial_engine.py`, `data/map_config.json`

The map management system translates **CS2 world coordinates** (typical values: -2000 to +2000 on X/Y) into **pixel coordinates** on the map texture (0.0 to 1.0 normalized), and vice versa.



MapMetadata (immutable dataclass for each map):

Field	Example (Dust2)	Purpose
<code>pos_x</code>	-2476	X of top-left corner in world units
<code>pos_y</code>	3239	Y of top-left corner in world units
<code>scale</code>	4.4	Game units per texture pixel
<code>z_cutoff</code>	None	Level separator (multi-level maps only)
<code>level</code>	"single"	Type: "single", "upper", "lower"
<code>texture_width</code>	1024	Radar texture width in pixels
<code>texture_height</code>	1024	Radar texture height in pixels

MapManager (`core/map_manager.py`):

The MapManager is the high-level component that coordinates map loading for the interface:

Method	Purpose
<code>get_map_metadata(map_name)</code>	Returns MapMetadata for the requested map
<code>load_radar_texture(map_name)</code>	Loads PNG texture of the radar map from cache or disk
<code>get_available_maps()</code>	List of supported maps with metadata
<code>world_to_pixel(pos, map_name)</code>	Shortcut for world → pixel coordinate conversion

Supported multi-level maps:

Map	z_cutoff	Levels	Notes
Nuke	-495	upper / lower	Two plant sites on different floors
Vertigo	11700	upper / lower	Skyscraper with two playable areas
All others	None	single	Single-level maps

SpatialEngine (`core/spatial_engine.py`):

The SpatialEngine adds **spatial reasoning** capabilities to the base coordinate system:

Method	Input	Output	Use
<code>distance_2d(pos_a, pos_b)</code>	Two world coordinates	Float (world units)	Engagement distance computation
<code>is_visible(pos_a, pos_b, obstacles)</code>	Two pos + obstacles	Bool	Line of sight (simplified)
<code>get_zone(pos, map_name)</code>	Coordinate + map	String (e.g., "A_site")	Tactical zone classification
<code>nearest_cover(pos, map_name)</code>	Coordinate + map	Cover coordinate	Positional suggestion

AssetManager (`core/asset_manager.py`):

Handles loading of map textures and UI assets:


Asset	Format	Typical size	Cache
Map texture (radar)	PNG 1024×1024	~500KB	Yes (in-memory)
Player icons	PNG 32×32	~2KB	Yes
Fonts (JetBrains Mono, YUPIX)	TTF	~200KB	Yes (registered in Kivy)
Theme wallpapers	PNG 1920×1080	~2MB	Lazy load

12.12 Observability and Logging

File: `Programma_CS2_RENAN/observability/logger_setup.py` **Related files:**

`backend/storage/state_manager.py` , `backend/services/telemetry_client.py`

The observability system ensures that every significant event is **traceable, structured, and persistent**. The telemetry client (`telemetry_client.py`) uses `httpx` (asynchronous HTTP) for non-blocking sending of metrics and events, avoiding network latency affecting application performance.



Daemons monitored by StateManager:

Daemon	Field in CoachState	Typical values
Scanner	hltv_status	"Scanning", "Paused", "Error"
Digester (worker)	ingest_status	"Processing", "Idle", "Error"
Teacher (trainer)	ml_status	"Learning", "Idle", "Error"
Global	status	"Paused", "Training", "Idle", "Error"

Sentry Integration (`observability/sentry_setup.py`):

The system includes optional integration with **Sentry** for remote error tracking, with a **double opt-in** approach for privacy:

Feature	Implementation
Double opt-in	The user must (1) set <code>SENTRY_DSN</code> as an environment variable AND (2) activate the flag in settings
PII Scrubbing	All personal data (player names, Steam IDs, paths) are removed before sending
Breadcrumb sanitization	Navigation breadcrumbs are cleaned of sensitive information
Non-blocking	If Sentry is not configured, the system continues normally (Phase 4 of startup)
Enriched context	Every error includes: app version, daemon state, demo count, operating system

Logger Setup (`observability/logger_setup.py`):

The centralized logging system provides structured logs with:


Feature	Detail
Format	<code>`TIMESTAMP</code>
File handler	<code>logs/cs2_analyzer.log</code> (append mode, automatic rotation)
Console handler	Only <code>WARNING+</code> to avoid polluting output
Naming convention	<code>get_logger("cs2analyzer.<module>")</code> — hierarchical namespace
Levels used	<code>DEBUG</code> (development), <code>INFO</code> (normal operations), <code>WARNING</code> (non-critical anomalies), <code>ERROR</code> (failures), <code>CRITICAL</code> (termination)

Training metrics exposed in CoachState: `current_epoch` , `total_epochs` , `train_loss` , `val_loss` , `eta_seconds` , `belief_confidence` , `system_load_cpu` , `system_load_mem` .

12.13 Reporting and Visualization

File: `Programma_CS2_RENAN/reporting/visualizer.py` , `report_generator.py` **Related files:** `backend/processing/heatmap_engine.py`

The reporting system transforms raw data into **understandable visualizations** for the user.



HeatmapEngine — Technical details:

Feature	Detail
Thread-safety	<code>generate_heatmap_data()</code> runs in a separate thread (does not block UI)
Texture creation	Only in the main thread (Kivy OpenGL requirement)
Type	2D Gaussian occupancy (parametric blur kernel)
Differential	Subtracts pro heatmap from user heatmap → red (too much time), blue (too little)
Hotspot	Identifies position clusters for positional training

MatchVisualizer (`reporting/visualizer.py`) — specialized rendering methods:

The MatchVisualizer extends the reporting capability with 6 high-quality rendering methods (cf. section 12.26 for algorithmic details):

Method	Input	Output	Description
<code>render_skill_radar(user, pro)</code>	Metrics dict	PNG radar chart	5-axis pentagon with pro baseline overlay
<code>render_trend_chart(history, metric)</code>	Historical list	PNG line chart	Temporal trend with moving average
<code>render_heatmap(positions, map_name)</code>	Tick coordinates	PNG heatmap	2D Gaussian on map texture
<code>render_differential(user_pos, pro_pos)</code>	Two position sets	PNG diff heatmap	Red (excess) / Blue (deficit) / Green (alignment)
<code>render_critical_moments(events)</code>	Round events list	PNG timeline	Color markers for kill/death/bomb
<code>render_comparison_table(user, pro)</code>	Two stats profiles	HTML/PNG table	Percentage delta for each metric

ReportGenerator (`reporting/report_generator.py`):


Assembles all visual elements into a **complete PDF document**:

1. **Header**: Player name, date, demo analyzed
2. **Executive Summary**: Key metrics (KPR, ADR, KAST, Rating)
3. **Radar Chart**: 5-axis comparison with pro baseline
4. **Trend Charts**: Trends of the last N matches
5. **Differential Heatmap**: Where to position better
6. **Critical Moments**: The 5 decisive moments of the match
7. **Coach's Advice**: Top-3 prioritized insights
8. **Footer**: Generated by Macena CS2 Analyzer v.X.X

12.14 Quota and Limit Management

The system implements a **monthly quota** mechanism to prevent abuse of processing resources and guarantee fair load distribution.

Limit	Value	Enforcement
<code>MAX_DEMOS_PER_MONTH</code>	10	<code>StorageManager.can_user_upload()</code>
<code>MAX_TOTAL_DEMOS</code>	100	<code>StorageManager.can_user_upload()</code>
<code>MIN_DEMOS_FOR_COACHING</code>	10	Threshold for full personalized coaching
Monthly reset	Automatic	<code>PlayerProfile.last_upload_month</code> vs current month



12.15 Fault Tolerance and Recovery

The system is designed to **never lose data** and to **recover automatically** from almost all types of failure.

FAULT TOLERANCE MECHANISMS

Zombie Task Cleanup
On startup: 'processing' tasks
→ reset to 'queued'
Automatic recovery without loss

Automatic Backup
On startup: 'startup_auto' checkpoint
Rotation: 7 daily + 4 weekly
Full database copy

Service Supervisor
Monitors external services
Auto-restart with exponential backoff
Max 3 attempts/hour

HLTV Circuit Breaker
MAX_FAILURES=10, RESET=3600s
State:
CLOSED → OPEN → HALF_OPEN
Prevents cascade failure

Connection Pooling
20 persistent connections
30s timeout for WAL contention
Automatic health checks


Graceful Degradation
Coaching: 4 fallback levels
GhostEngine: (0,0) on error
Every service has a plan B

Failure and recovery matrix:

Failure scenario	Recovery mechanism	Data loss
Application crash during parsing	Zombie Task Cleanup on restart	None — task restarted
Database corruption	Restore from most recent automatic backup	Up to 24 hours of data
HLTV service unreachable	Circuit Breaker <code>_CircuitBreaker</code> (MAX_FAILURES=10, RESET_WINDOW_S=3600): after 10 consecutive failures the circuit opens and blocks requests for 1 hour, then passes to HALF_OPEN for a test. Prevents cascade failure on external APIs.	None — pro data delayed
ML model not loadable	Fallback to random weights (GhostEngine) or base coaching	None — degraded quality
Insufficient RAM during training	Automatic early stopping, checkpoint saved	None — last valid checkpoint
Disk full	Database Governor detects and notifies via ServiceNotification	Prevention — no data written
Corrupt/truncated demo	IntegrityChecker rejects before parsing	None — demo ignored with warning
Ollama unavailable	LLM fallback to pure RAG, then base coaching	None — degraded narrative quality
Steam/FACEIT API timeout	Retry with exponential backoff (max 3 attempts)	None — profile not updated
Playwright browser crash	BrowserManager: recreates instance, retries operation	None — HLTV scraping delayed
WAL conflict (lock contention)	30s timeout with automatic retry	None — operation delayed
Corrupt ML checkpoint	Fallback to previous checkpoint (versioning)	Partial — loses last training
Zero-norm query vector (M-07)	<code>VectorIndex.search()</code> returns <code>None</code> with warning — RAG pipeline continues without crash	None — empty RAG result
Corrupt settings.json file (M-08)	<code>load_user_settings()</code> validates structure, falls back to default if non-dict	None — default settings
Logger not initialized (C-09)	Early bootstrap uses <code>print()</code> before logger init	None — messages on stdout

Failure scenario	Recovery mechanism	Data loss
DB handle leak on integrity check (H-02)	<code>_verify_integrity()</code> uses <code>try/finally</code> + <code>PRAGMA quick_check</code>	None — connection always closed
Masked batch errors in RAP training (H-03)	Removed <code>except (KeyError, TypeError): continue</code> — <code>train_step()</code> requires dict with <code>'loss'</code> key	None — errors now visible


Graceful degradation — The Fallback Chain:



12.16 Complete User Journey — 4 Main Flows

This section describes the **4 main flows** a user goes through while using Macena CS2 Analyzer.


Flow 1: Demo Upload and Analysis



Detail: Internal CoachingService Pipeline

When `generate_new_insights()` is invoked, the CoachingService internally runs:





▼


8. Persistence
CoachingInsight in DB
(5-20 per match)

LessonGenerator — generation of structured lessons from demo:


Parameter	Value	Meaning
ADR_STRONG	75	$ADR \geq 75 \rightarrow$ "Good average damage"
HS_STRONG	0.40	$HS\% \geq 40\% \rightarrow$ "Good headshot accuracy"
KAST_STRONG	0.70	$KAST \geq 70\% \rightarrow$ "Good round contribution"
Output format	Section 1 (Overview) + Section 2 (Strengths) + Section 3 (Areas for improvement) + Section 4 (Pro tip)	Standard lesson structure

The LessonGenerator operates in dual mode: if Ollama is available, it generates natural language lessons via `LLMService.generate_lesson()` . If not available, it uses structured templates with numeric data inserted into preformed phrases.

Flow 2: Tactical Visualization



Flow 3: Automatic ML Retraining




Flow 4: AI Chat with the Coach

CoachingDialogueEngine — Multi-Turn Pipeline:


The AI dialogue follows a 5-phase pipeline for each user message:

Phase	Component	Description
1	Intent Classification	Classifies the question: coaching, stats query, comparison, general
2	Context Retrieval	Retrieves the last N matches, recent insights, player profile
3	RAG Augmentation	Searches the knowledge base for the 3 most relevant documents (similarity search 384-dim)
4	COPER Augmentation	If available, adds past coaching experiences with positive outcome
5	Response Generation	Generates response via Ollama (if available) or structured RAG template

Parameter	Value	Purpose
MAX_CONTEXT_TURNS	6	Maximum number of turns kept in context
MAX_RESPONSE_WORDS	100	Word limit for LLM response
SIMILARITY_THRESHOLD	0.5	Minimum threshold for RAG results
Tone	"coaching, encouraging"	System prompt for positive tone



Flow 5: Diagnostics and Maintenance




Tool matrix by scenario:

Scenario	Recommended tool	Time
Quick check post-deploy	headless_validator.py	~10s
Suspicious database	db_inspector.py	~30s
Demo not parsed	demo_inspector.py	~5s
Training diverges	Ultimate_ML_Coach_Debugger.py	~2min
Complete check-up	Goliath_Hospital.py	~3min
Deep ML audit	brain_verify.py	~5min

12.17 Tool Suite — Validation and Diagnostics (tools/)

Directory: tools/ (root) + Programma_CS2_RENAN/tools/ **Main files:** headless_validator.py , db_inspector.py , demo_inspector.py , brain_verify.py (planned, not yet implemented), Goliath_Hospital.py , Ultimate_ML_Coach_Debugger.py , _infra.py

The tool suite is a **collection of 41 scripts** distributed across two directories (tools/ root with 29 scripts and Programma_CS2_RENAN/tools/ with 12 scripts) that form a multi-level validation pyramid. Each tool has a precise purpose and can be executed independently, but together they form a quality assurance system covering every aspect of the project.



Headless Validator (tools/headless_validator.py) — the mandatory regression gate (Dev Rule 9). Executed after **every** development task with 23 phases and 319 checks:

Phase	Verification	Detail
1	Environment	Python $\geq$ 3.10, critical dependencies present (torch, kivy, sqlalchemy, demoparser2)
2	Core Import	<code>config.py</code> , <code>spatial_data.py</code> , <code>lifecycle.py</code> — the fundamental modules load without errors
3	Backend Import	<code>nn/</code> , <code>processing/</code> , <code>storage/</code> , <code>services/</code> , <code>coaching/</code> — all backend subsystems importable
4	DB Schema	The 21 SQLAlchemy tables are created correctly, the relationships are valid
5	Configuration	<code>METADATA_DIM</code> , paths, constants — consistent and reachable values
6	ML Smoke	Model instantiation (JEPA, RAP, MoE) with random weights — verify dimensions and forward pass
7	Observability	<code>get_logger()</code> working, <code>StateManager</code> initializable, writable log path

Shared Infrastructure (`tools/_infra.py`):

Component	Role
<code>BaseValidator</code>	Base class for all validators — unified console output, pass/fail/warning count
<code>path_stabilize()</code>	<code>sys.path</code> normalization to avoid inconsistent relative imports
<code>Console</code> (Rich)	Color output with tables, progress bars, status emoji
<code>Severity</code> enum	4 levels: <code>INFO</code> , <code>WARNING</code> , <code>ERROR</code> , <code>CRITICAL</code>

DB Inspector (`tools/db_inspector.py`) — deep database inspection:

- Opens `database.db` and `hltv_metadata.db` separately
- For each table: counts records, verifies schema, checks index integrity
- Displays space metrics (file size, WAL pages, fragmentation)
- Detects anomalies: unexpected empty tables, orphan records, out-of-range timestamps
- Output: color Rich table with status for each table

Demo Inspector (`tools/demo_inspector.py`) — demo file inspection:

- Verifies magic bytes (`PBDEMS2\0` for CS2, `HL2DEMO\0` for legacy)
- Checks file size (1KB min, 5GB max)
- Extracts header metadata without full parsing
- Detects partial corruption (truncation, damaged header)
- Output: report on demo validity with error details if invalid

Brain Verify (`tools/brain_verify.py` + `brain_verification/`, 16 sections) — *(planned, not yet implemented)*:

The intelligence verification system is designed with **16 thematic sections** that will cover **118 rules** of quality:



BRAIN VERIFY — 16 SECTIONS (118 RULES)

§1 Dimensional Contracts
METADATA_DIM, latent_dim,
input shapes

§2 Forward Pass
Smoke test all models
with synthetic input

§3 Loss Functions
Gradient flow, NaN check,
loss monotonicity


§4 Training Pipeline
DataLoader, optimizer,
scheduler config

§5 Feature Engineering
Vectorizer alignment,
normalization bounds

§6 Coaching Pipeline
COPER → Hybrid → RAG
fallback chain

§7 Data Integrity
DB schema, FK constraints,
orphan detection

§8 Concept Alignment
16 coaching concepts,
label correctness



Goliath Hospital (`tools/Goliath_Hospital.py`) — multi-department diagnostics:

The "Goliath Hospital" is the **most complete diagnostic system** of the project, organized as a hospital with 10 specialized departments:

Department	Name	Checks
1	Emergency Room (ER)	Critical imports, immediate crashes, path resolution
2	Radiology	File/directory structure, missing files, permissions
3	Pathology	DB schema, data integrity, anomalous records
4	Cardiology	Session Engine, daemon heartbeat, IPC
5	Neurology	ML models, forward pass, gradient flow
6	Oncology	Dead code, unused imports, circular dependencies
7	Pediatrics	Onboarding, first-user flows, wizard
8	Intensive Care (ICU)	Concurrency, race conditions, WAL contention
9	Pharmacy	Dependencies, versions, compatibility
10	Tool Clinic	Validation of other tools (meta-test)

Ultimate ML Coach Debugger (`tools/Ultimate_ML_Coach_Debugger.py`) — neural belief falsification:

This tool performs a **3-phase audit** on the ML pipeline:

- 1. **Structural Verification** — Tensor sizes, model architecture, learnable parameters
- 2. **Behavioral Verification** — Forward pass with real data, gradient flow, loss convergence
- 3. **Belief Falsification** — Tests whether the model "believes" wrong things: overconfident predictions, systematic biases, degenerate patterns


Other tools:

Tool	File	Purpose
<code>backend_validator.py</code>	<code>tools/</code>	Verifies backend import and interface consistency
<code>build_tools.py</code>	<code>tools/</code>	PyInstaller build automation
<code>user_tools.py</code>	<code>tools/</code>	Utilities for the end user (reset, export)
<code>dev_health.py</code>	<code>tools/</code>	Quick development health check
<code>context_gatherer.py</code>	<code>tools/</code>	Gathers context for debug reports
<code>dead_code_detector.py</code>	<code>tools/</code>	Identifies unreferenced dead code
<code>sync_integrity_manifest.py</code>	<code>tools/</code>	Updates <code>integrity_manifest.json</code> with SHA-256 hashes
<code>project_snapshot.py</code>	<code>tools/</code>	Complete snapshot of project state
<code>ui_diagnostic.py</code>	<code>tools/</code>	UI-specific diagnostics (Qt/Kivy)
<code>build_pipeline.py</code>	Root <code>tools/</code>	Complete build pipeline
<code>Feature_Audit.py</code>	Root	Audit of the 25-dim feature vector
<code>Sanitize_Project.py</code>	Root	Cleanup of temporary files, cache, artifacts
<code>audit_binaries.py</code>	Root	Verifies integrity of executables and .pt

12.18 Test Suite Architecture (`tests/`)

Directory: `Programma_CS2_RENAN/tests/` **Total files:** 94 test files + `conftest.py` + 10 forensics scripts + 15 verification scripts

The test suite is organized according to the **test pyramid principle**: many unit tests (fast, isolated), fewer integration tests (slower, with real dependencies), and few end-to-end tests (complete but costly).



Conftest (`tests/conftest.py`) — shared fixtures:

Fixture	Scope	Description
<code>test_db</code>	function	In-memory SQLite database with all tables created, auto-cleanup
<code>sample_match_stats</code>	function	<code>PlayerMatchStats</code> with realistic data derived from real matches
<code>mock_demo_data</code>	function	Mock <code>DemoData</code> with frames and events for parsing tests
<code>tmp_models_dir</code>	function	Temporary directory for <code>.pt</code> checkpoints
<code>coaching_service</code>	function	<code>CoachingService</code> configured with test DB

Automated Suite (`tests/automated_suite/`):

Category	File	Coverage
Smoke	<code>test_smoke_*.py</code>	Critical imports, model instantiation, DB connection
Unit	<code>test_unit_*.py</code>	Pure functions, calculations, transformations
Functional	<code>test_functional_*.py</code>	Complete pipelines, workflows
E2E	<code>test_e2e_*.py</code>	Demo upload → parsing → coaching → insight
Regression	<code>test_regression_*.py</code>	Specific bug fixes (G-01 label leakage, G-07 Bayesian, etc.)

Test per subsystem:

Subsystem	Test files	What they test
NN Core	<code>test_jepa_model.py</code> , <code>test_rap_model.py</code> , <code>test_advanced_nn.py</code>	Forward pass, tensor dimensions, gradient flow
Coaching	<code>test_coaching_service.py</code> , <code>test_hybrid_engine.py</code> , <code>test_correction_engine.py</code>	Coaching pipeline, insight prioritization, fallback chain
Processing	<code>test_vectorizer.py</code> , <code>test_heatmap.py</code> , <code>test_feature_eng.py</code>	Feature extraction, normalization, METADATA_DIM=25
Storage	<code>test_database.py</code> , <code>test_models.py</code> , <code>test_backup.py</code>	CRUD, schema, backup/restore, WAL mode
Data Sources	<code>test_demo_parser.py</code> , <code>test_hltv_scraper.py</code> , <code>test_steam_api.py</code>	Parsing, scraping, API timeout/retry
Analysis	<code>test_game_theory.py</code> , <code>test_belief.py</code> , <code>test_role_engine.py</code>	Analysis engines, calibration, heuristics
Knowledge	<code>test_rag.py</code> , <code>test_experience_bank.py</code> , <code>test_knowledge_graph.py</code>	Retrieval, COPER effectiveness, KG query
Ingestion	<code>test_ingest_pipeline.py</code> , <code>test_registry.py</code>	Complete pipeline, deduplication

Forensics (`tests/forensics/` , 10 scripts):

Forensics scripts are diagnostic tools for post-mortem investigations:


Script	Purpose
diagnose_training_failure.py	Analyzes training logs to identify divergence, NaN, gradient explosion
inspect_model_weights.py	Weight distribution, layer statistics, dead neurons
replay_ingestion.py	Re-runs ingestion of a specific demo with verbose logging
trace_coaching_path.py	Traces the path of an insight from demo to UI
db_consistency_check.py	Verifies consistency across the 3 databases

Verification Scripts (15 files in tests/ root):

One-shot verification scripts to validate specific aspects of the system:

Script	Verification
verify_feature_pipeline.py	METADATA_DIM=25 respected in all paths
verify_training_cycle.py	4 training phases complete without error
verify_db_schema.py	21 tables present with correct schema
verify_coaching_pipeline.py	Demo → insight end-to-end path
verify_imports.py	All modules importable without circular errors
verify_rag_index.py	Knowledge base indexed with correct dimensions (384-dim)
verify_pro_baseline.py	Pro baseline loaded and valid
verify_hltv_sync.py	HLTV sync service configured correctly

Test strategy by criticality:




Identified coverage gaps:

The following modules have high complexity (>500 LOC) but limited test coverage:

Module	LOC	Test coverage	Risk
training_orchestrator.py	733	Low	HIGH — multi-phase orchestration
session_engine.py	538	Medium	MEDIUM — testable only with subprocess
coaching_service.py	585	Medium	MEDIUM — 4 fallback levels
experience_bank.py	751	Low	HIGH — complex COPER logic
tensor_factory.py	686	Low	HIGH — DataLoader/Dataset
coach_manager.py	878	Low	HIGH — global training state


12.19 The 12 Phases of Systematic Remediation

The project went through a **12-phase remediation** process that overall resolved **370+ issues** identified during progressive quality audits. Each phase focused on a specific category of problems, from critical bug fixes to architectural restructuring.



Detail per phase:

Phase	Focus	Issues resolved	Key example
1	Import and base structure	~20	Circular imports, path resolution, missing <code>__init__.py</code>
2	Security and secrets	~25	Hard-coded API keys → env vars/keyring, input validation
3	Data pipeline and features	~30	G-01 label leakage, G-02 normalization bounds, feature alignment
4	Database and schema	~30	WAL mode enforcement, missing indexes, schema migration safety
5	Dead code and cleanup	~25	G-06 elimination of <code>nn/advanced/</code> , unused imports, duplicate files
6	Analysis engines	~30	Graceful degradation for all 10 engines, edge case handling
7	ML pipeline	~35	G-07 Bayesian calibration, gradient clipping, checkpoint versioning
8	Coaching and COPER	~30	G-08 experience decay, RAG index validation, fallback chain
9	UI and UX	~25	F8-XX visual feedback, state consistency, error prevention
10	Resilience and concurrency	~40	F5-35 threading, Event, timeout enforcement, circuit breaker
11	Observability and logging	~35	F5-33 structured logging, correlation IDs, Sentry integration
12	Deep debug and wiring	~45	18 issues: wiring verification, integration testing, edge case audit



Directory `reports/` :

Each remediation phase produced a detailed report saved in the `reports/` directory. Each report includes:

- Numbered list of issues found with code (G-XX or F-XX)
- Problem description with affected files and lines
- Implemented solution with conceptual diff
- Verification: confirmation that the fix did not introduce regressions

Key fixes by code (G-XX):

Code	Phase	Description	Impact
G-01	3	Label leakage in <code>ConceptLabeler.label_tick()</code> — future features leaked into training labels	CRITICAL — invalidated VL-JEPA training
G-02	4	Danger zone in <code>vectorizer.py</code> — normalization without bounds check	HIGH — NaN propagation in training
G-06	5	Dead code in <code>backend/nn/advanced/</code> — 3 unreferenced files	MEDIUM — confusion and accidental imports
G-07	7	Bayesian death estimator not calibrated — static model without auto-calibration	HIGH — inaccurate predictions
G-08	8	COPER experience bank without decay — stale experiences never removed	MEDIUM — coaching based on outdated data

Fixes per feature (F-XX):

Code	Area	Description
F3-29	Processing	ResNet resize inconsistent — image dimensions not normalized
F5-19	Services	Matplotlib rendering without error handling — crash on empty stats
F5-21	Storage	Explicit <code>session.commit()</code> inside context manager — double commit
F5-22	Security	Hard-coded API keys in source — migrated to env vars + keyring
F5-33	Observability	<code>print()</code> debugging — migrated to structured logging
F5-35	Control	<code>time.sleep()</code> loop for stop — migrated to <code>threading.Event</code>
F6-XX	Analysis	Analysis engines without graceful degradation — crash on incomplete input
F7-XX	Knowledge	RAG without index validation — inconsistent embedding dimensions
F8-XX	UI	Kivy widget without visual feedback — silent actions confuse the user

12.20 Pre-commit Hooks and Quality Gates

File: `.pre-commit-config.yaml`


The project uses a **pre-commit hooks** system that activates automatically before every commit, preventing non-compliant code from reaching the repository.

4 Local Hooks (project custom):

Hook	Script	Timeout	Description
headless-validator	tools/headless_validator.py	20s	23 phases, 319 checks regression gate — the most important
dead-code-detector	tools/dead_code_detector.py	15s	Identifies unreferenced imports and functions
integrity-manifest-check	tools/sync_integrity_manifest.py	10s	Verifies SHA-256 manifest hash consistency
dev-health-quick	tools/dev_health.py	10s	Quick project health check

Standard Hooks (7 from `pre-commit/pre-commit-hooks` + 2 Python):

Hook	Source	Description
trailing-whitespace	pre-commit/pre-commit-hooks	Removal of trailing whitespace
end-of-file-fixer	pre-commit/pre-commit-hooks	Required final newline
check-yaml	pre-commit/pre-commit-hooks	YAML syntax validation
check-json	pre-commit/pre-commit-hooks	JSON syntax validation
check-added-large-files	pre-commit/pre-commit-hooks	Blocks files >1MB (prevents accidental <code>.pt</code>)
check-merge-conflict	pre-commit/pre-commit-hooks	Detects unresolved merge conflict markers
detect-private-key	pre-commit/pre-commit-hooks	Blocks commits of private keys
black	psf/black	Automatic Python formatting (line length 100, target py3.12)
isort	pycqa/isort	Automatic import sorting (black profile, line length 100)



12.21 Build, Packaging, and Deployment

Key files: `packaging/windows_installer.iss` , `setup_new_pc.bat` , `export_env.bat`

The project includes a build and packaging system for distributing the desktop application on Windows.

3 Requirements Files:

File	Purpose	Dependencies
<code>requirements.txt</code>	Base — main dependencies for execution	28 packages (torch, kivy, sqlmodel, demoparser2, playwright, httpx, etc.)
<code>requirements-ci.txt</code>	CI/CD — adds test and analysis tools	pytest, coverage, mypy, black, isort, pre-commit
<code>requirements-lock.txt</code>	Lock — exact versions for reproducible builds	Exact pin of each dependency and sub-dependency

Windows Packaging (`packaging/windows_installer.iss`):

- Inno Setup installer with graphical wizard
- Custom CS2 Analyzer icon
- Desktop and Start menu shortcut
- Estimated size: ~500MB (includes PyTorch and models)

- Target: Windows 10/11 64-bit

Setup Scripts:

Script	Purpose
<code>setup_new_pc.bat</code>	Environment setup from scratch: Python, pip, venv, dependencies
<code>export_env.bat</code>	Export of required environment variables
<code>run_full_training_cycle.py</code>	Full ML training cycle orchestration

12.22 Alembic Migration System

Directory: `alembic/` (root) + `Programma_CS2_RENAN/backend/storage/migrations/`

The project uses **two separate Alembic setups** to manage database schema migrations in an organized way.

Alembic Root (`alembic/` , 13 versions):

Contains the main schema migrations, from initial table creation to the most recent changes. Each migration:

- Has a unique hash as identifier
- Contains `upgrade()` (applies change) and `downgrade()` (reverts change)
- Is idempotent — can be reapplied without errors
- Is tested on a production-like schema before deploy

Backend Migrations (`backend/storage/migrations/` , 2 versions):

Migrations specific to the storage layer, kept separate for modularity. They handle changes to tables that are the exclusive responsibility of the `backend/storage/` module.

Automatic execution: Migrations are executed automatically during phase 3 of the startup sequence (`main.py`), ensuring the schema is always up to date before any DB operation.

12.23 Main Ingestion Orchestrator (`run_ingestion.py`)

File: `Programma_CS2_RENAN/run_ingestion.py`

`run_ingestion.py` is the **orchestrating heart** of the entire ingestion pipeline. It is the largest file dedicated to ingestion and coordinates all phases from demo discovery to persisting results in the database.

12+ main functions:

Function	Role
<code>discover_demos()</code>	Scans configured directories, filters valid <code>.dem</code> files
<code>validate_demo_file()</code>	Checks magic bytes, size, pre-parsing integrity
<code>process_single_demo()</code>	Complete pipeline for a demo: parse → extract → enrich → persist
<code>batch_process()</code>	Processes multiple demos in sequence with progress tracking
<code>enrich_round_stats()</code>	Post-processing: computes trade kills, blind kills, flash assists
<code>compute_hltv_rating()</code>	Computes HLTV 2.0 rating for each player
<code>assign_dataset_split()</code>	Assigns 70% train / 15% val / 15% test with temporal split
<code>generate_coaching_insights()</code>	Invokes CoachingService to generate 5-20 insights
<code>update_ingestion_task()</code>	Updates task status in DB (queued → processing → completed/failed)
<code>cleanup_failed_tasks()</code>	Cleans up failed tasks, resets to queued if recoverable
<code>report_progress()</code>	Structured logging of overall progress
<code>handle_duplicate_detection()</code>	Verifies via Registry whether the demo has already been processed

ResourceManager (`ingestion/resource_manager.py`):


The ResourceManager manages **hardware resources** during ingestion to avoid system overload:

Parameter	Default Value	Purpose
CPU threshold	80%	Pauses ingestion if CPU exceeds this threshold
RAM threshold	85%	Pauses ingestion if RAM exceeds this threshold
Disk check	Yes	Verifies sufficient disk space before processing
Throttle delay	2s	Delay between consecutive tasks for cooldown

12.24 HLTV Sync Service and Background Daemon

File: `Programma_CS2_RENAN/hltv_sync_service.py` **Related files:** `backend/data_sources/hltv/` , `backend/services/telemetry_client.py`

The HLTV Sync Service is a **background daemon** that automatically synchronizes pro player data from HLTV.org. It operates as a service monitored by the Console's `ServiceSupervisor` .




Component	File	Role
CircuitBreaker	<code>circuit_breaker.py</code>	Protects against cascade failure: CLOSED → OPEN (after 10 fails) → HALF_OPEN (test) → CLOSED
BrowserManager	<code>browser_manager.py</code>	Headless Playwright browser management with FlareSolverr for Cloudflare bypass
CacheProxy	<code>cache_proxy.py</code>	Local cache of HLTV responses to reduce requests
RateLimiter	<code>rate_limiter.py</code>	Configurable delay between requests to respect ToS

HLTV Internal Architecture — Collectors and Selectors:

The `ingestion/hltv/` tree also contains specialized modules for data collection:

Module	Type	Description
<code>player_collector.py</code>	Collector	Collects pro player profiles from HLTV pages
<code>team_collector.py</code>	Collector	Collects rosters and team statistics
<code>match_collector.py</code>	Collector	Collects match results and demo links
<code>stat_selector.py</code>	Selector	Extracts specific statistics from parsed HTML pages
<code>demo_selector.py</code>	Selector	Identifies and downloads demo links from match pages

Detailed operating cycle:




12.25 RASP Guard — Runtime Code Integrity

File: `Programma_CS2_RENAN/observability/rasp.py` **Data file:** `data/integrity_manifest.json`

The RASP (Runtime Application Self-Protection) Guard is the **first check** executed on application startup (Phase 1 of the boot sequence). It verifies that no source file has been modified compared to the integrity manifest.

RASP Guard Components:

Component	Description
<code>RASPGuard</code>	Main class — loads the manifest, verifies hashes, raises exception
<code>integrity_manifest.json</code>	JSON file with SHA-256 hashes of each critical source file
<code>IntegrityError</code>	Custom exception — immediate termination with detailed log
<code>sync_integrity_manifest.py</code>	Tool to update the manifest after legitimate changes

Verification flow:**12.26 MatchVisualizer — Advanced Rendering**


File: `Programma_CS2_RENAN/reporting/visualizer.py`

The `MatchVisualizer` extends the reporting system (section 12.13) with **6 specialized rendering methods** for generating charts and visualizations.

6 Rendering methods:

Method	Output	Description
<code>render_skill_radar()</code>	PNG radar chart	5-axis pentagon (Mechanics, Positioning, Utility, Timing, Decision) with pro overlay
<code>render_trend_chart()</code>	PNG line chart	Historical trend of key metrics (KPR, ADR, KAST) over N matches
<code>render_heatmap()</code>	PNG heatmap	2D Gaussian occupancy on tactical map
<code>render_differential_heatmap()</code>	PNG diff heatmap	user-pro subtraction: red=too much, blue=too little, green=optimal
<code>render_critical_moments()</code>	PNG annotated timeline	Timeline with color markers for key moments (kills, deaths, plants)
<code>render_comparison_table()</code>	PNG/HTML table	user vs pro comparison table with percentage delta for each metric

Differential heatmap — algorithm:



12.27 Runtime Data and Configuration Files

The project includes several data files that are read or written during execution. These files are not Python code but are essential for the system to function.

data/map_config.json — CS2 map configuration:

Contains the coordinate transformation parameters for each supported map (Dust2, Mirage, Inferno, Nuke, Vertigo, Overpass, Ancient, Anubis):

```

{
  "de_dust2": {
    "pos_x": -2476, "pos_y": 3239, "scale": 4.4,
    "z_cutoff": null, "level": "single"
  },
  "de_nuke": {
    "pos_x": -3453, "pos_y": 2887, "scale": 7.0,
    "z_cutoff": -495, "level": "multi"
  }
}

```

data/integrity_manifest.json — RASP integrity manifest:

Dictionary `{file_path: sha256_hash}` for all critical Python files. Generated/updated by `tools/sync_integrity_manifest.py`. Verified on startup by `RASPGuard` (section 12.25).

data/training_progress.json — Training progress:

Persists training state between restarts: current epoch, best loss, learning rate, last processed demo. Consulted by the Teacher daemon to resume training from the exact point where it stopped.

user_settings.json — User settings:

Level 2 configuration file (cf. section 12.3). Saved in the project directory, contains customizable preferences: demo paths, UI theme, language, feature flags.

12.28 Root-Level Entry Points


Besides `main.py`, the project includes several **executable scripts** at the root level that serve as alternative entry points for specific operations:

Script	Lines	Purpose	Invocation
<code>console.py</code>	~61KB	Complete interactive TUI console with Rich — command registry, CLI and interactive modes, complete system management	<code>python console.py</code>
<code>run_ingestion.py</code>	1,057	Standalone ingestion orchestrator (cf. 12.23) — batch demo processing	<code>python run_ingestion.py [path]</code>
<code>goliath.py</code>	~200	Launcher for Goliath Hospital diagnostics (cf. 12.17)	<code>python goliath.py</code>
<code>schema.py</code>	~100	Generation and visualization of the current DB schema	<code>python schema.py</code>
<code>run_full_training_cycle.py</code>	~150	Standalone 4-phase complete training (JEPA → Pro → User → RAP)	<code>python run_full_training_cycle.py</code>
<code>hflayers.py</code>	~50	Utility for Hopfield layers (NCPs integration)	Import only

Interactive console (`console.py`, ~61KB):

The console is the **most powerful entry point** for advanced operators. It offers a Rich-based TUI (Terminal User Interface) with:


- **Command Registry:** System of registered commands with autocompletion
- **CLI Mode:** Single-shot command execution (`console.py status`, `console.py train`)
- **Interactive Mode:** REPL shell with custom prompt and command history
- **Rich Panels:** Dashboard with color tables, progress bars, status tree view




Architectural Summary

Macena CS2 Analyzer is a **layered and modular** application organized across 5 architectural layers with 3 separate processes.








Complete technology stack (28 dependencies):

Category	Library	Version	Role in the project
ML Core	PyTorch	2.x	Neural networks (JEPA, RAP, MoE, NeuralRoleHead, WinProb)
ML Ext	ncps	latest	Liquid Time-Constant Networks (LTC) for RAP memory
ML Ext	hflayers	latest	Hopfield layers for associative attention
UI	Kivy	2.x	Cross-platform UI framework
UI	KivyMD	1.x	Material Design widgets
DB	SQLModel	latest	ORM (Pydantic + SQLAlchemy)
DB	SQLAlchemy	2.x	Underlying database engine
DB	Alembic	1.x	Schema migrations
HTTP	requests	2.x	Steam/FACEIT API (synchronous)
HTTP	httpx	latest	Telemetry (asynchronous)
Scraping	playwright	latest	Headless browser for HLTV
Parsing	demoparser2	latest	CS2 demo parser (Rust-based, fast)
Data	pandas	2.x	Tabular data manipulation
Data	numpy	1.x	Numerical computations
Viz	matplotlib	3.x	Charts and visualizations
NLP	sentence-transformers	latest	384-dim embeddings for RAG
TUI	rich	latest	Color console, tables, progress bars
Logging	logging (stdlib)	—	Structured logging
Security	keyring	latest	Secure management of API keys
Monitoring	sentry-sdk	latest	Remote error tracking (optional)
Testing	pytest	latest	Test framework
Formatting	black	latest	Code formatting
Import	isort	latest	Import sorting
QA	pre-commit	latest	Pre-commit hooks
Crypto	hashlib (stdlib)	—	SHA-256 for RASP manifest
Concurrency	threading (stdlib)	—	Daemon threads, MLControlContext
IPC	subprocess (stdlib)	—	Session Engine as subprocess
Config	json (stdlib)	—	user_settings.json, map_config.json

The 3 application processes:

Process	Type	Responsibility	Communication
Main	Kivy GUI	User interface, rendering, interaction	DB polling every 10-15s
Daemon	Subprocess (Session Engine)	Scanner, Digester, Teacher, Pulse	stdin pipe (IPC) + shared DB
Optional Services	External processes	HLTV sync, local Ollama LLM	HTTP/API + Console supervision

Main data flows:



Remediation Note — Eliminated Code (G-06)

During the 12-phase remediation process (370+ issues resolved overall — cf. section 12.19), the `backend/nn/advanced/` directory was **completely eliminated** because it contained unreferenced dead code:

- `superposition_net.py` — An experimental superposition network never integrated into training or inference flow. Active superposition functionality is implemented in `layers/superposition.py` (used by the RAP Strategy layer).
- `brain_bridge.py` — An experimental model bridge never called by any module.
- `feature_engineering.py` — Duplicate feature engineering; the canonical version resides in `backend/processing/feature_engineering/`.

Rationale (G-06): Keeping dead code creates risks of confusion (which `superposition` is the real one?), increases maintenance surface, and can introduce accidental imports. The elimination was verified through static dependency analysis: no file in the project imported from `backend/nn/advanced/`.

Shared Utility — `round_utils.py`

File: `Programma_CS2_RENAN/backend/knowledge/round_utils.py`

The function `infer_round_phase(equipment_value)` is a **shared utility** used by both the coaching service and the knowledge system to classify the economic phase of a round. It resides in `backend/knowledge/` and is imported by modules in `services/` and `processing/`.

Equipment value	Returned phase
< \$1,500	"pistol"
\$1,500 – \$2,999	"eco"
\$3,000 – \$3,999	"force"
≥ \$4,000	"full_buy"

Architectural Strengths

1. **25-dimensional unified feature contract** — `METADATA_DIM = 25` enforces system-wide training/inference parity.
2. **3-level maturity gating** — Prevents premature model deployment with insufficient data.
3. **4-level coaching fallback** — COPER → Hybrid → RAG → Base guarantees insights are always delivered.
4. **Multi-model diversity** — JEPA, VL-JEPA, LSTM+MoE, RAP, and NeuralRoleHead contribute complementary inductive biases.
5. **Temporal split** — Prevents data leakage by ensuring chronological ordering.
6. **COPER feedback loop** — EMA-based effectiveness monitoring with decay of stale experience.
7. **Phase 6 analysis suite** — 10 analysis engines (role, win probability, game tree, belief, deception, momentum, entropy, blind spots, utility and economy, engagement distance).
8. **Threshold persistence** — Role thresholds survive restarts via the DB table `RoleThresholdRecord`.
9. **Configurable heuristics** — `HeuristicConfig` externalizes normalization bounds to JSON.
10. **LLM polishing** — Optional Ollama integration for natural language coaching narratives.
11. **Training Observatory** — 4-level introspection (Callback, TensorBoard, Maturity State Machine, Embedding Projector) with zero impact when disabled and callbacks isolated from errors.
12. **Neural Role Consensus** — Dual heuristic + NeuralRoleHead MLP classification with cold-start protection, guaranteeing reliable role assignment even with partial data.
13. **Per-Round Statistical Isolation** — The `RoundStats` model prevents contamination across rounds, enabling granular round-level coaching and per-round HLTV 2.0 evaluations.
14. **Quad-Daemon Architecture** — Complete separation between GUI and heavy work, with coordinated shutdown and automatic zombie task cleanup.
15. **Pervasive graceful degradation** — Every component has a fallback plan: the system never crashes, it always degrades in a controlled way.
16. **Three-Tier Storage Architecture** — Separation of `database.db` (core + knowledge, 18 tables), `hlvtv_metadata.db` (pro data, 3 tables), and `match_data/{id}.db` (per-match telemetry) to eliminate WAL contention and prevent uncontrolled monolith growth.
17. **Live Bayesian Calibration (G-07)** — The death estimator self-calibrates with `extract_death_events_from_db()` → `auto_calibrate()`, transforming from a static model into an adaptive system.

18. **Live Training Control (MLControlContext)** — Real-time pause/resume/stop/throttle of training via `threading.Event`, with custom exception `TrainingStopRequested` instead of `StopIteration`.
19. **Resilient Circuit Breaker** — `_CircuitBreaker` for external APIs (HLTV) with MAX_FAILURES=10, RESET_WINDOW_S=3600, prevents cascade failure with CLOSED → OPEN → HALF_OPEN pattern.
20. **5-Level Validation Pyramid** — Headless Validator → pytest → Backend Validator → Goliath Hospital → Brain Verify: each level progressively deeper, from quick smoke test (10s) to complete audit of 118 rules.
21. **RASP Guard** — Runtime Application Self-Protection with SHA-256 manifest: verifies source code integrity on startup, prevents accidental and intentional tampering.
22. **13-Hook Pre-commit Gate** — 4 custom hooks + 7 standard + 2 Python prevent non-compliant code from reaching the repository. Formatting, validation, dead code, integrity, merge conflict, and private key verified automatically.
23. **Hardware-Aware ResourceManager** — Automatic CPU/RAM throttling during ingestion: the system autonomously slows down if hardware resources are under pressure, without human intervention.
24. **Test Forensics** — 10 post-mortem investigation scripts to diagnose training failure, weight anomalies, coaching path traces, and DB consistency — the system's "black box".



ARCHITECTURAL STRENGTHS — THE 24 PILLARS

1. Unified 25-dim contract -
Everyone speaks the same language

2. 3-level maturity gate - No
premature release

3. 4-level coaching fallback - Never
empty-handed

4. Multi-model diversity - 5 brains > 1
brain

5. Temporal split - No time-travel
cheating

6. COPER feedback loop - Learns
from its own advice

7. Phase 6 Analysis (10 eng.) - 10
specialized detectives

8. Threshold persistence - Survives
restarts

9. Configurable heuristics - JSON
override

10. LLM polish (Ollama) - Advice sounds natural

11. Training Observatory - Report card for the brain

12. Neural Role Consensus - Two teachers compare notes

13. Per-Round isolation - Evaluates every question, not just the test

14. Quad-Daemon architecture - Responsive GUI, heavy work in background

15. Graceful degradation - The system never crashes


16. Three-Tier Storage - No contention between write and read

17. Live Bayesian calibration - Self-calibrates with the data

18. Live Training control - Pause/Stop without loss

19. Circuit Breaker - External API

...




End of document — Complete guide to Macena CS2 Analyzer

Total `.py` files in the project: **~406** (in `Programma_CS2_RENAN/`, ~480 project-wide) Total verified Python lines of code: **≈ 102,000+** AI subsystems covered: **8** (NN Core, VL-JEPA, RAP Coach, Coaching Services, Coaching Engines, Knowledge, Analysis, Processing + Training Observatory) Program subsystems covered: **18** (Startup, Lifecycle, Configuration, Session Engine, Desktop UI, Ingestion, Storage, Observability, Control Console, RASP Guard, HLTV Sync, Ingestion Orchestrator, ResourceManager, Tools Suite, Test Suite, Pre-commit, Build/Packaging, Alembic Migrations) Documented models: **6** (AdvancedCoachNN/TeacherRefinementNN, JEPA, VL-JEPA, RAPCoachModel, NeuralRoleHead, WinProbabilityNN) Documented analysis engines: **10** (Role, WinProb, GameTree, Belief, Deception, Momentum, Entropy, Blind Spots, Utility and Economy, Engagement Distance) Documented coaching engines: **7** (HybridEngine, CorrectionEngine, ExplainabilityGenerator, NNRefinement, ProBridge, TokenResolver, LongitudinalEngine) Additional services documented: **7** (CoachingDialogue, LessonGenerator, LLMService, VisualizationService, ProfileService, AnalysisService, TelemetryClient) Database tables documented: **21** (distributed across three-tier storage architecture: `database.db`, `hlvtv_metadata.db`, `match_data/{id}.db`) UI screens documented: **14** (Wizard, Home, Coach, Tactical

Viewer, Settings, Help, Match History, Match Detail, Performance, User Profile, Profile, Pro Comparison, Steam Config, FACEIT Config) — Qt/PySide6 (primary) + Kivy/KivyMD (legacy) Documented daemons: **4** (Scanner, Digester, Teacher, Pulse) Documented validation tools: **41** (Headless Validator, Brain Verify, Goliath Hospital, DB Inspector, Demo Inspector, ML Coach Debugger, Backend Validator, Dead Code Detector, Dev Health, Feature Audit, etc.) Documented test files: **94** (+ conftest.py, 10 forensics, 15 verification scripts) Headless validator phases: **23** (319 automated checks) Documented pre-commit hooks: **13** (4 custom local + 7 standard + 2 Python) Architectural pillars: **24** (including Three-Tier Storage, Live Bayesian Calibration, Live Training Control, Circuit Breaker, Validation Pyramid, RASP Guard, Pre-commit Gate, HW-aware ResourceManager, Forensics) Issues resolved through remediation: **368** (in 12 systematic phases) Documented remediation phases: **12** (with G-XX and F-XX codes)

Cross-Part Interconnection Map

The three parts of the documentation form an interconnected system. This map shows the main dependencies between sections:



Key cross-references:

From (Part)	To (Part)	Connection
Part 1 §VL-JEPA	Part 3 §10 (Training)	The two-stage protocol described in Part 1 is implemented by the training regime in Part 3
Part 1 §Data Sources	Part 3 §12.6 (Ingestion)	Data sources feed the ingestion pipeline
Part 2 §Coaching Engines	Part 3 §12.16 Flow 1	Coaching engines generate the insights displayed in the UI
Part 2 §Control Module	Part 3 §12.4 (Session Engine)	The control module manages the 4 daemons
Part 2 §Processing	Part 1 §NN Core	The vectorizer produces the METADATA_DIM=25 consumed by the models
Part 3 §9 (Database)	Part 2 §Knowledge	The RAG/COPER tables are used by the knowledge system
Part 3 §11 (Loss)	Part 1 §VL-JEPA	The documented loss functions implement VL-JEPA training
Part 3 §12.17 (Tools)	Part 3 §12.18 (Tests)	The validation pyramid includes both tools and test suite

Technical Glossary

Term	Definition
ADR	Average Damage per Round — average damage dealt per round
BCE	Binary Cross-Entropy — loss for binary classification
COPER	Coaching through Past Experiences and References — coaching based on past experiences
EMA	Exponential Moving Average — exponential moving average used for JEPA target encoder
HLTV 2.0	HLTV.org rating system based on KPR, survival, impact, damage
InfoNCE	Noise Contrastive Estimation — contrastive loss to align representations
JEPA	Joint-Embedding Predictive Architecture — architecture that predicts target embeddings from context
KAST	Kill/Assist/Survive/Trade — percentage of rounds with positive contribution
KG	Knowledge Graph — knowledge graph with entities and relationships
KPR	Kills Per Round — average kills per round
LTC	Liquid Time-Constant — neural networks with learned time constants
METADATA_DIM	Dimension of the unified feature vector (25 in the project)
MoE	Mixture of Experts — network with specialized experts and gating
MVVM	Model-View-ViewModel — architectural pattern for UI/logic separation
RAG	Retrieval-Augmented Generation — generation enriched by retrieval
RAP	Reasoning Action Planning — model for strategic coaching
RASP	Runtime Application Self-Protection — runtime application protection
TUI	Terminal User Interface — user interface in the terminal
VICReg	Variance-Invariance-Covariance Regularization — regularization for embedding diversity
VL-JEPA	Vision-Language JEPA — extension with linguistic concepts
WAL	Write-Ahead Logging — SQLite mode for concurrent reads/writes
Z-score	Number of standard deviations from the mean — measures how far a value is from the norm

Author: Renan Augusto Macena